


Benchmarking Optimization

LINFO2266

Pierre Schaus



Benchmarking Optimization





- You have a set of approaches or algorithms to solve an optimization problem.
- You have a set of benchmark instances (supposed to be representative of an unknown future instance).
- How do you compare those algorithms?



Evaluation through Benchmarking

- \mathcal{P} is the set of problems, \mathcal{S} is the set of solvers , $n_p = | \mathcal{P} |$, $n_s = | \mathcal{S} |$
- For each solver s , problem p , computing times = $t_{p,s}$



Instance	A	B	C
A	100	10	15
A	30	7	12
 A	40	45	35
 A	50	55	40
	10	30	15
	20	50	25

Which car is the best?

Objectives of evaluating different solvers/strategies

- Derive general conclusions provided the benchmarks are representative enough
- Quantify how much faster/slower the solver is compared to others
- Estimate the probability that a solver is the best on any unseen instance (requires statistical testing, not covered)

Proposal 1: The average

- Computing the average computing time over all instances?

Most difficult instances
dominates ✖



Discarding Unsolved
Instances



Bias towards most
robust solvers ✖

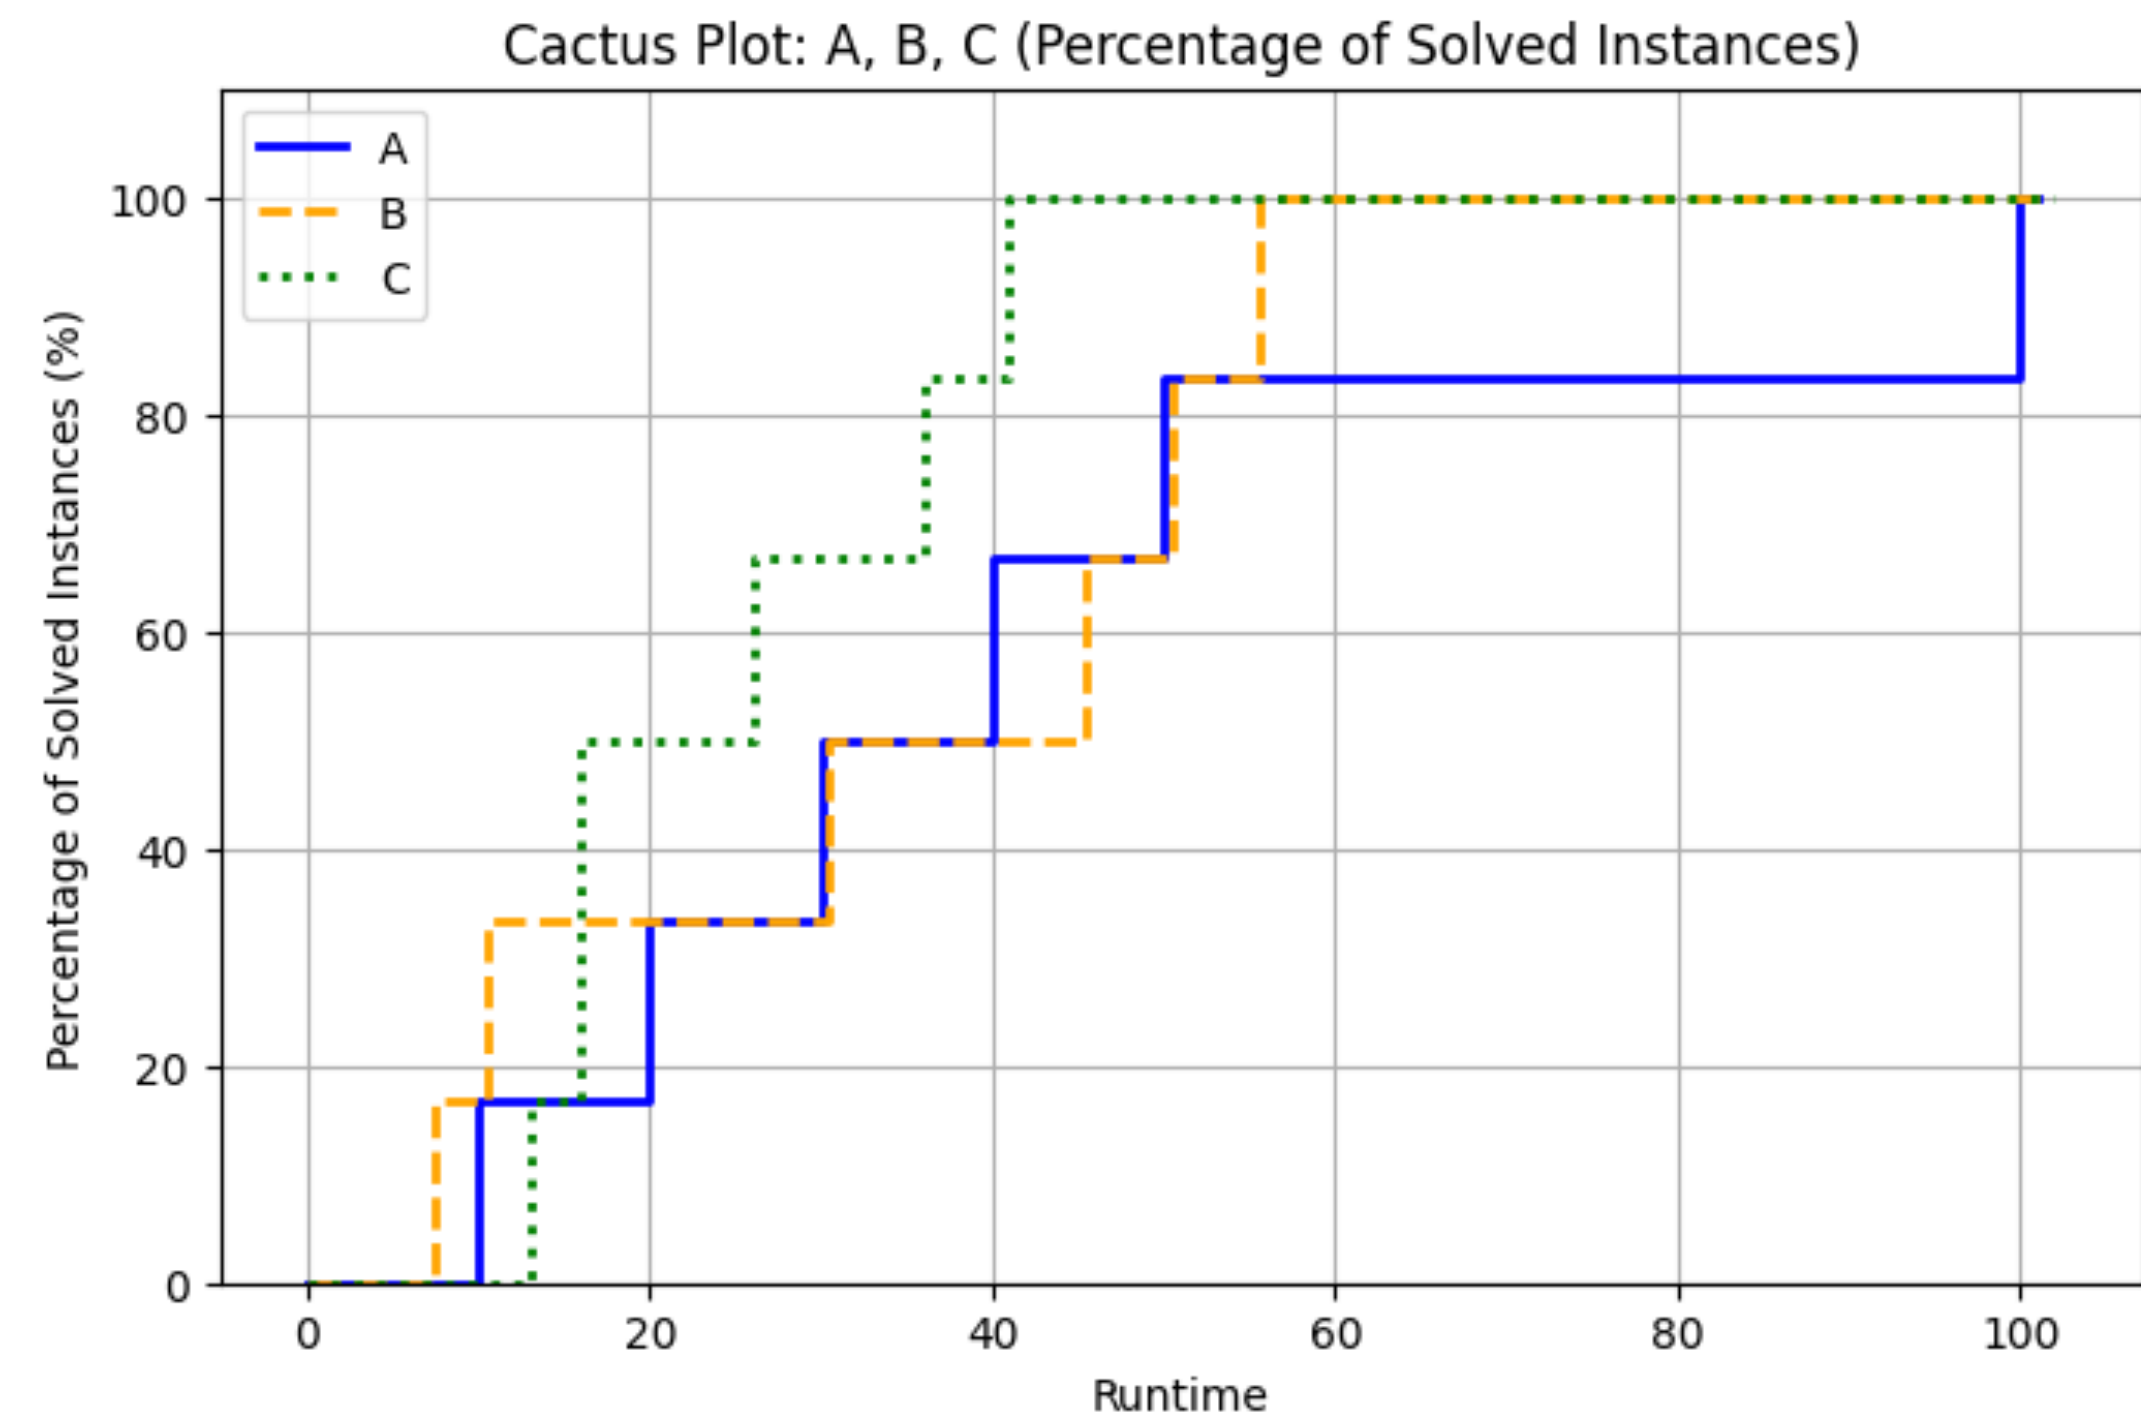
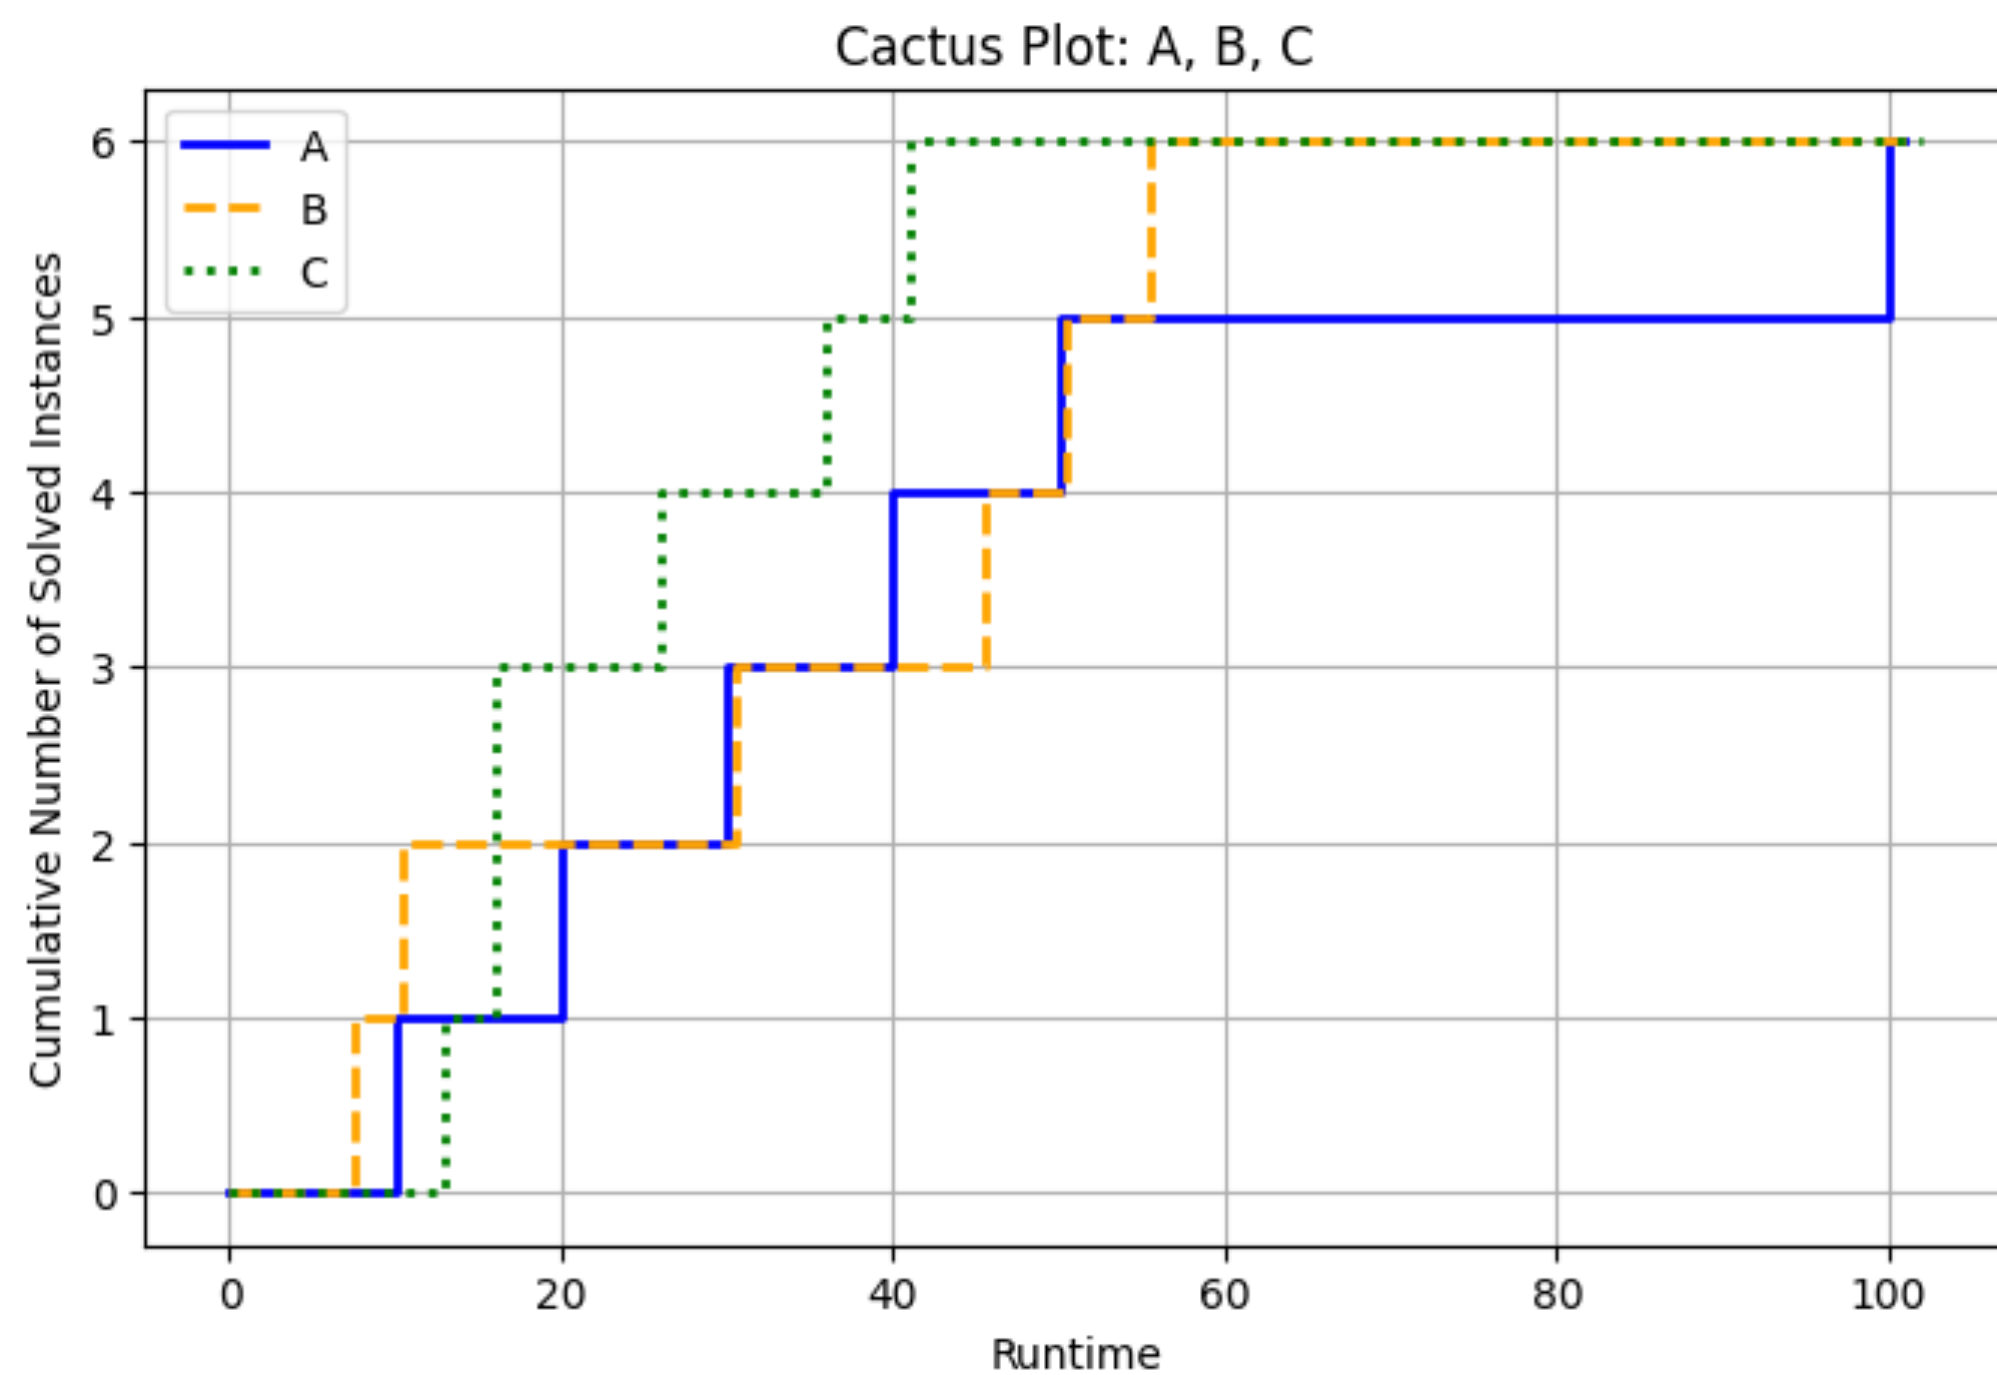
Proposal 2: Count solved instances

- What about ranking the solvers by number of solved instances?
- Very coarse grained ✖
- Relative order per instance is lost ✖









Proposal 3: Cactus Plot

- Fine grained 
- Relative order per instance is lost 

$$\kappa_s(\tau) = \frac{1}{n_p} \text{size}\{p \in \mathcal{P} : t_{p,s} \leq \tau\}$$



Proposal 4: ratio to *virtual best*

Instance	A	B	C	min	$r_{i,A}$	$r_{i,B}$	$r_{i,C}$
	100	10	15	10	10.0	1.0	1.5
	30	7	12	7	4.3	1.0	1.7
 	40	45	35	35	1.1	1.3	1.0
 	50	55	40	40	1.3	1.4	1.0
	10	30	15	10	1.0	3.0	1.5
	20	50	25	20	1.0	2.5	1.3

Virtual best

ratio to virtual best

►
$$r_{p,s} = \frac{t_{p,s}}{\min\{t_{p,s} : s \in \mathcal{S}\}}$$

Dolan, E. D., & Moré, J. J. (2002). **Benchmarking optimization software with performance profiles**. Mathematical programming, 91(2), 201-213.

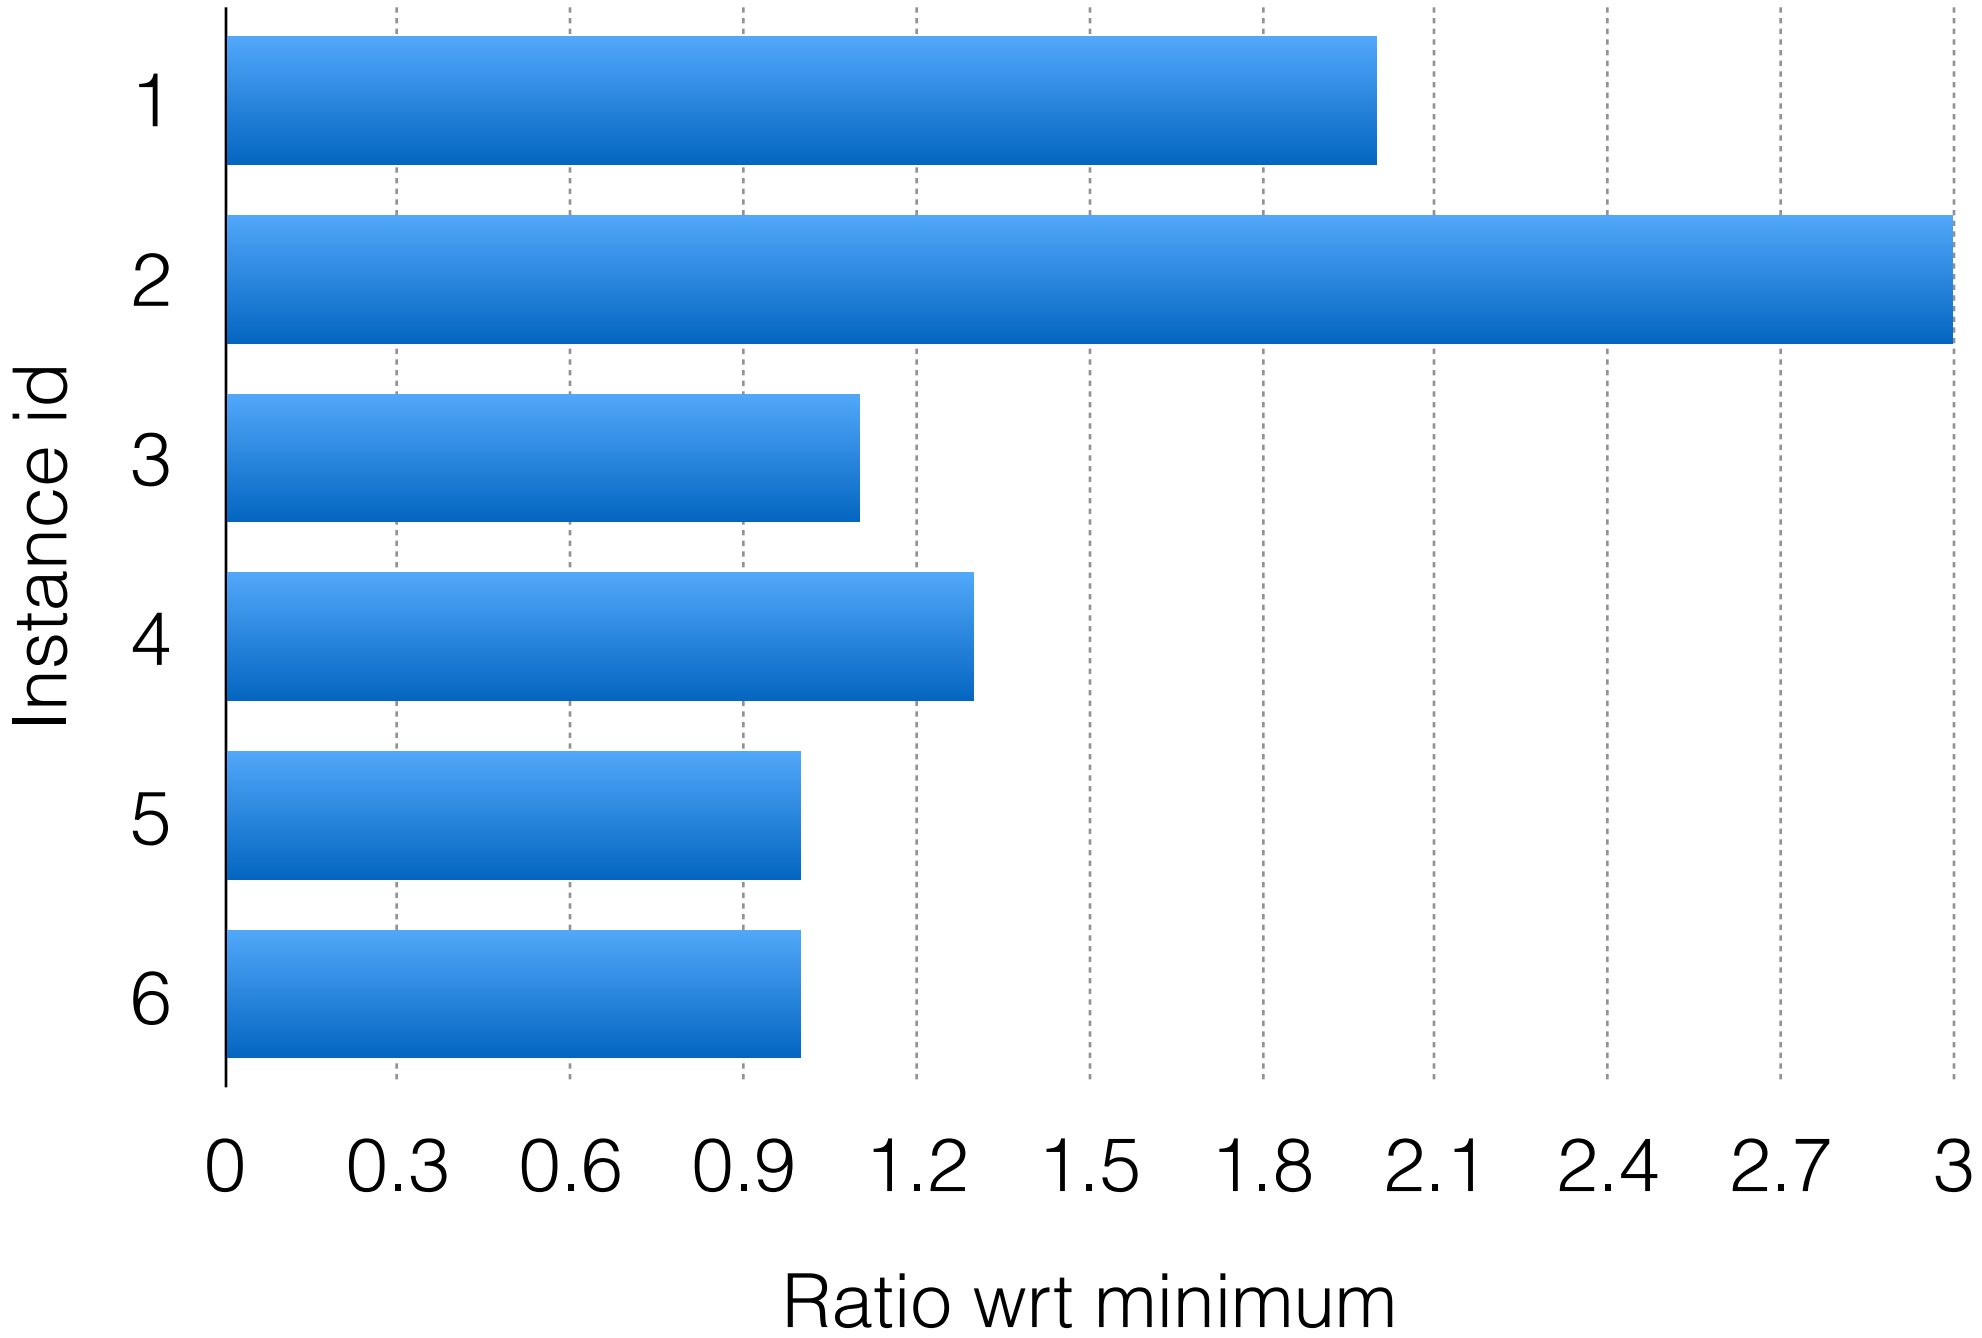
Bar-Plot of the ratio?



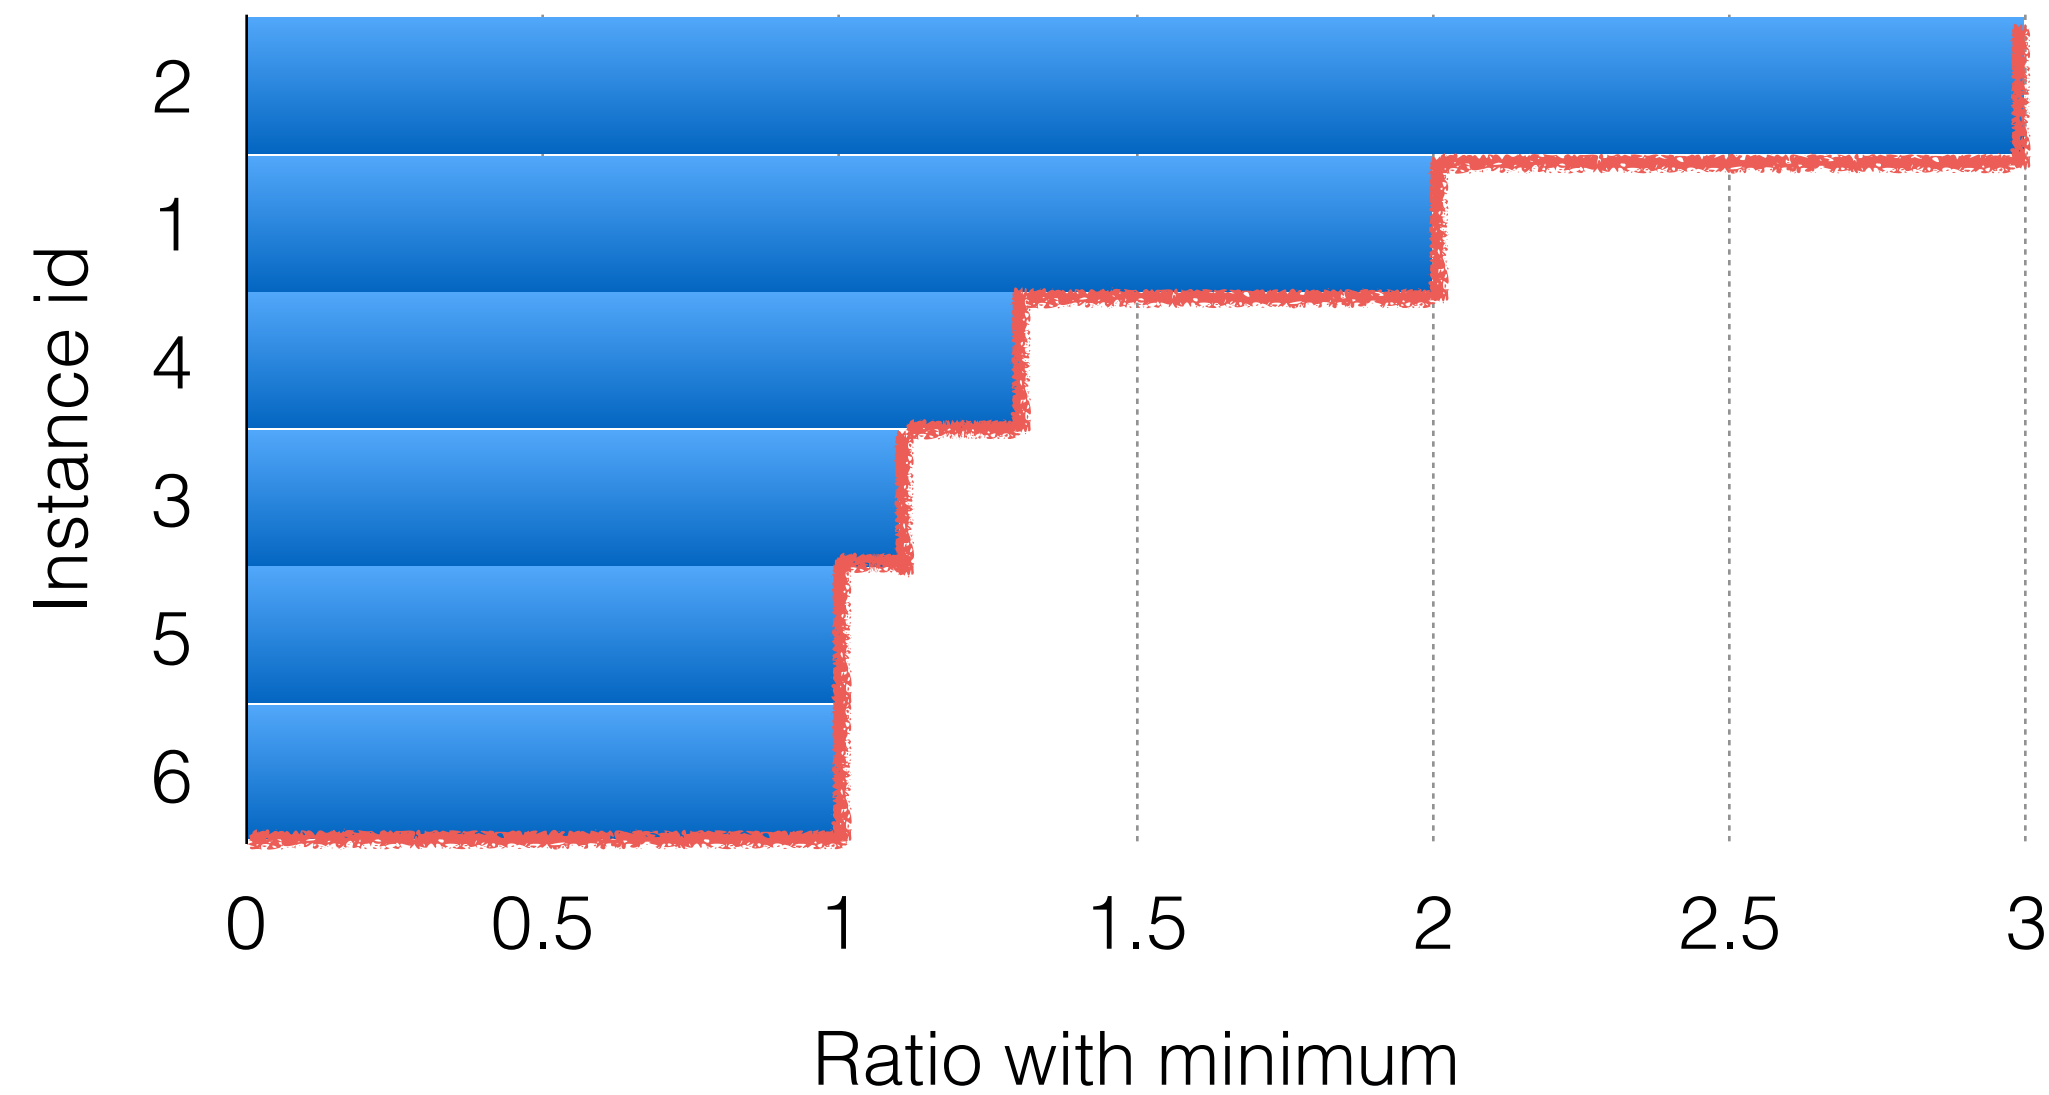
Solver A

Not very informative 🤔

Instance id	A	min	$r_{i,A}$
1	20	10	2.0
2	21	7	3.0
3	40	35	1.1
4	50	40	1.3
5	10	10	1.0
6	20	20	1.0

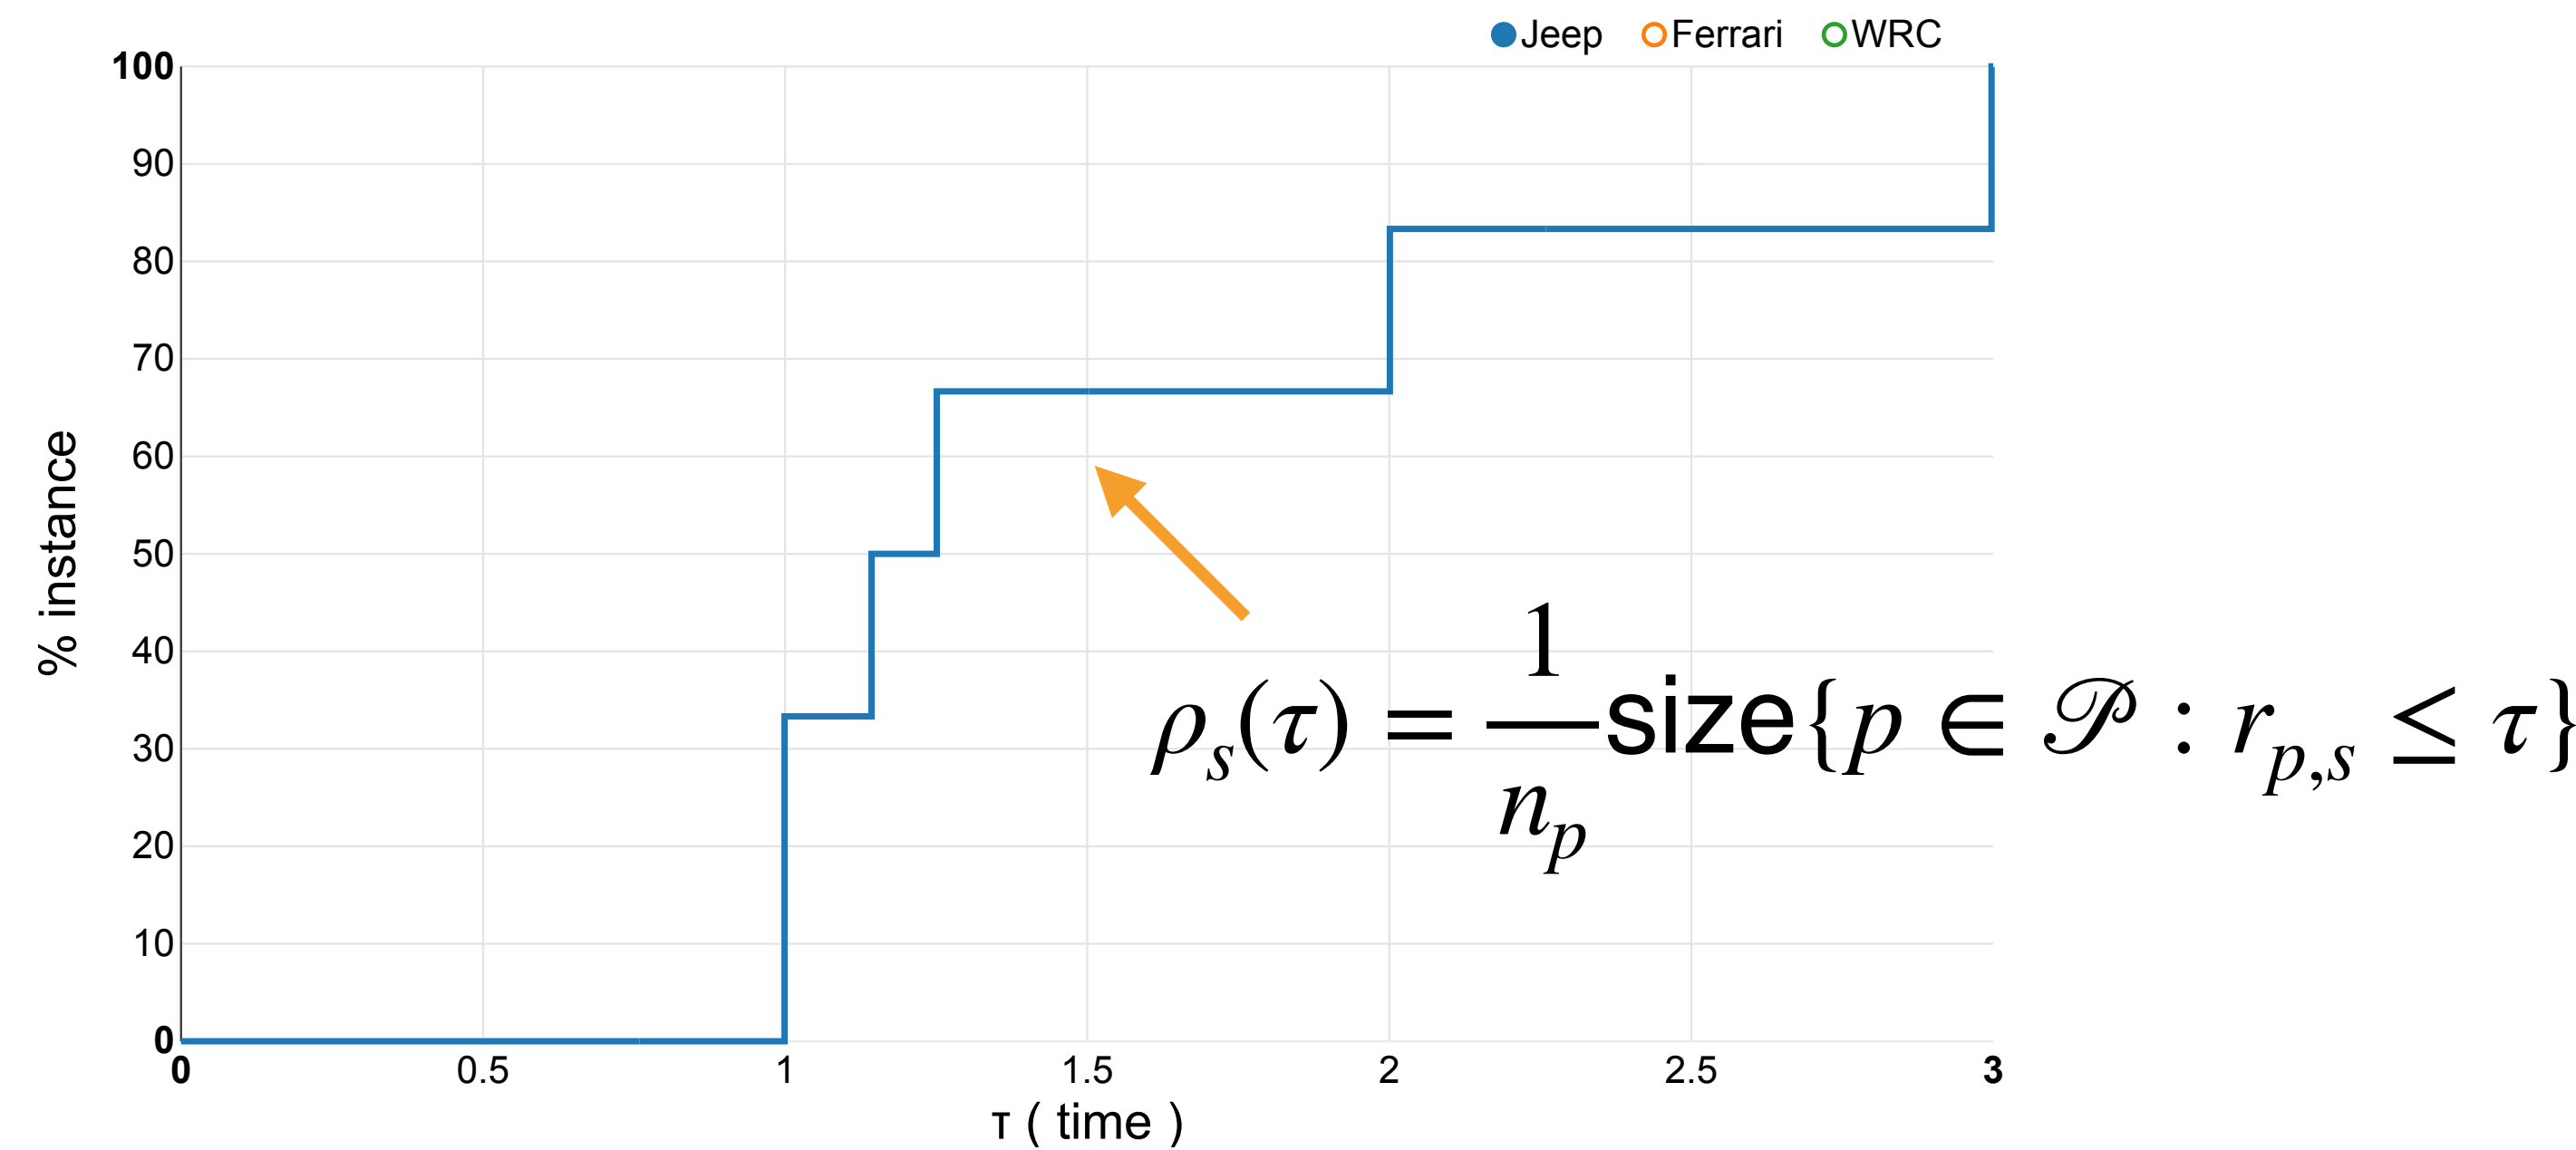


Performance Profile = Profile of virtual best ratios

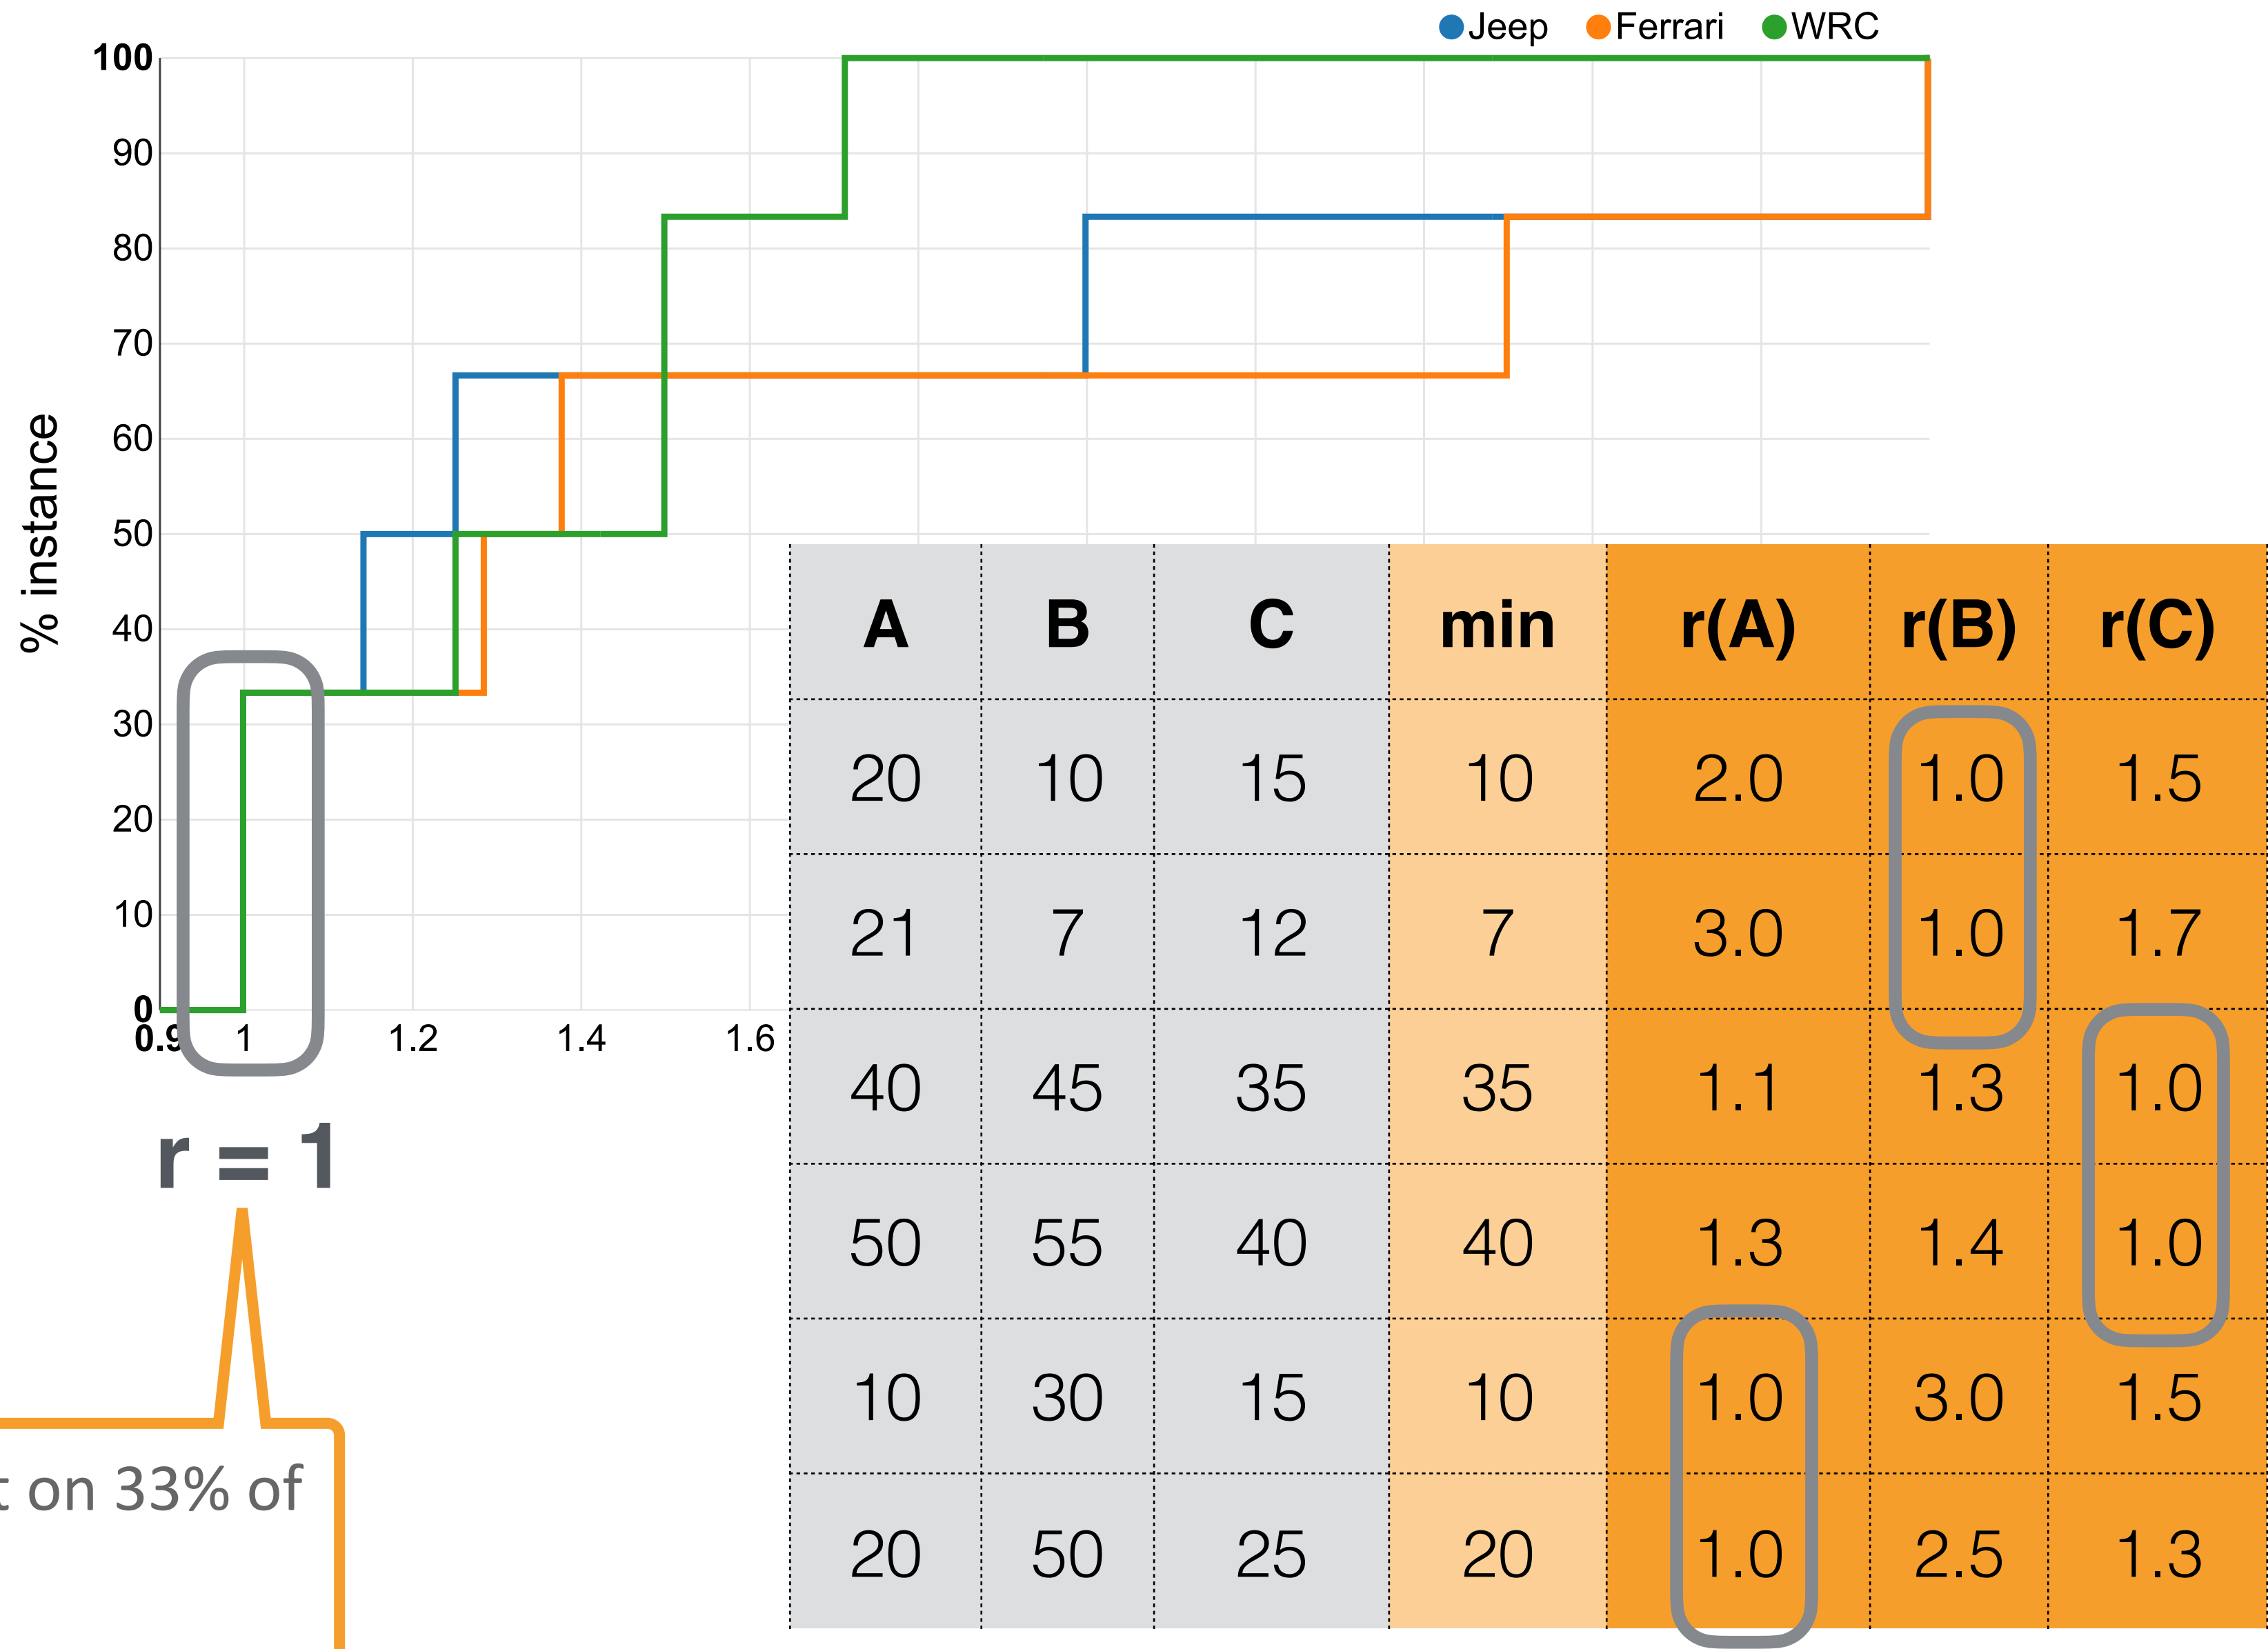


Instance id	A	min	$r_{i,A}$
2	21	7	3.0
1	20	10	2.0
4	50	40	1.3
3	40	35	1.1
5	10	10	1.0
6	20	20	1.0

$$r_{p,s} = \frac{t_{p,s}}{\min\{t_{p,s} : s \in \mathcal{S}\}}$$



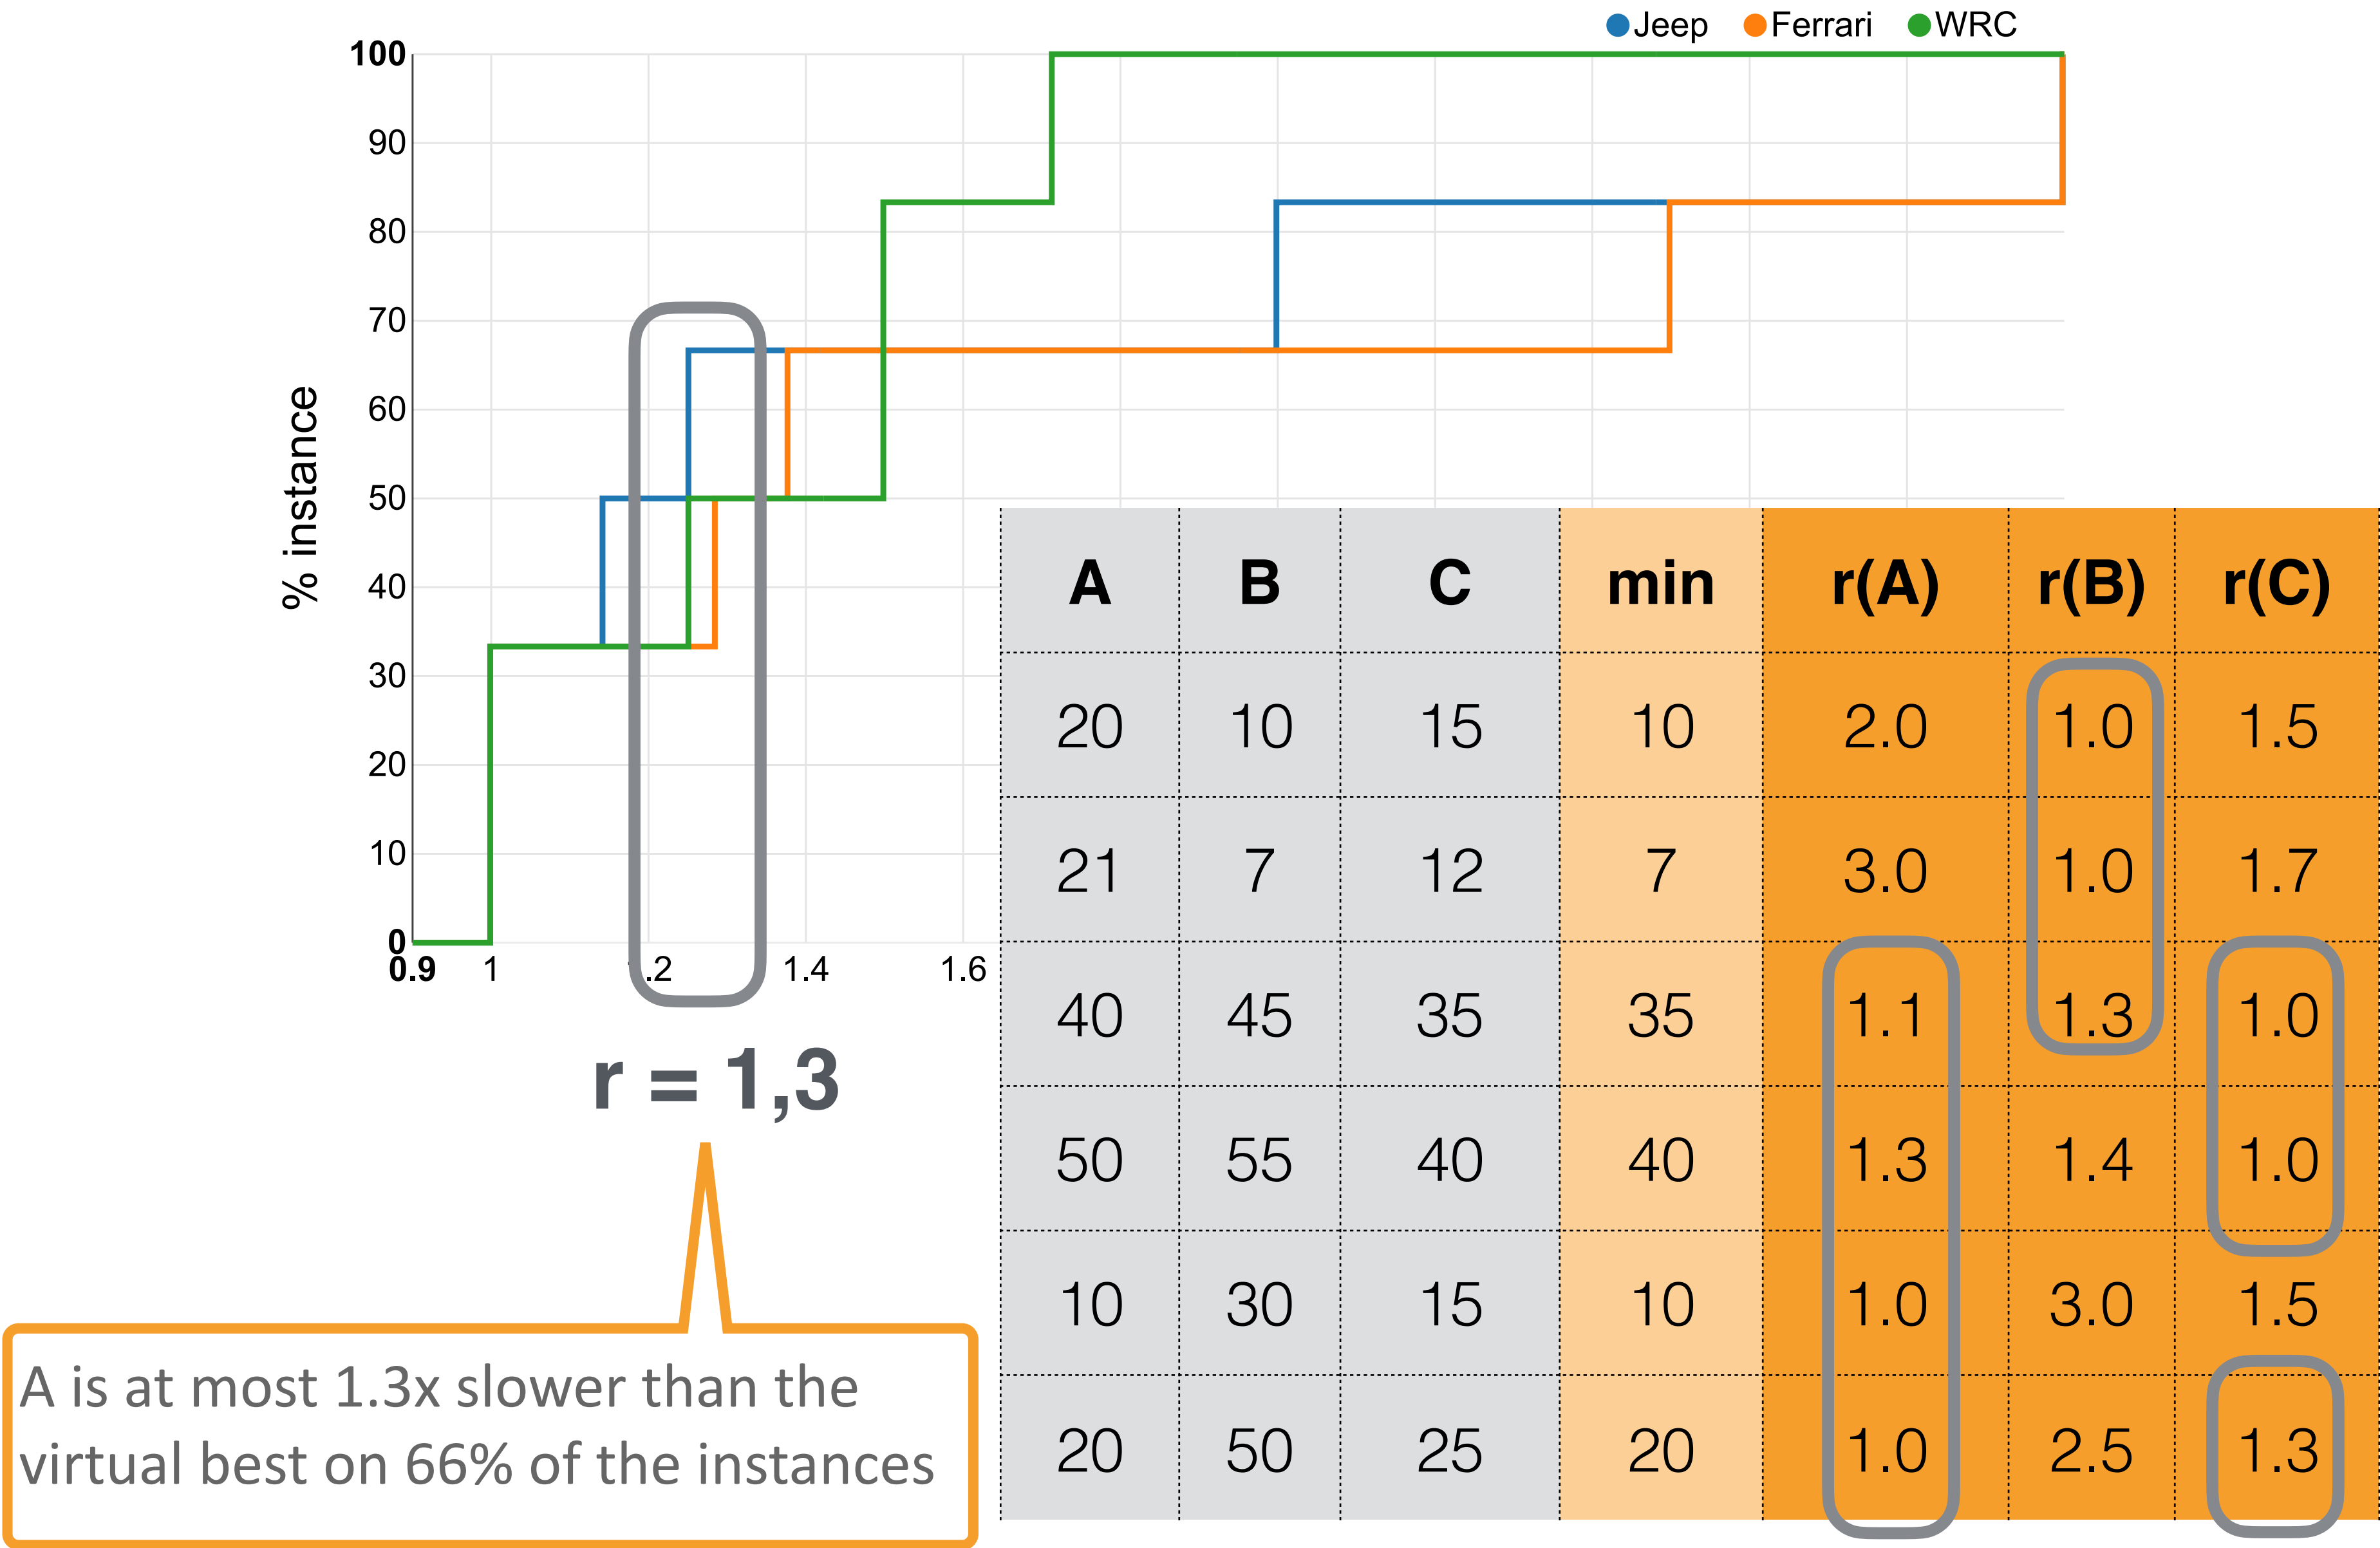
Performances Profiles Interpretation



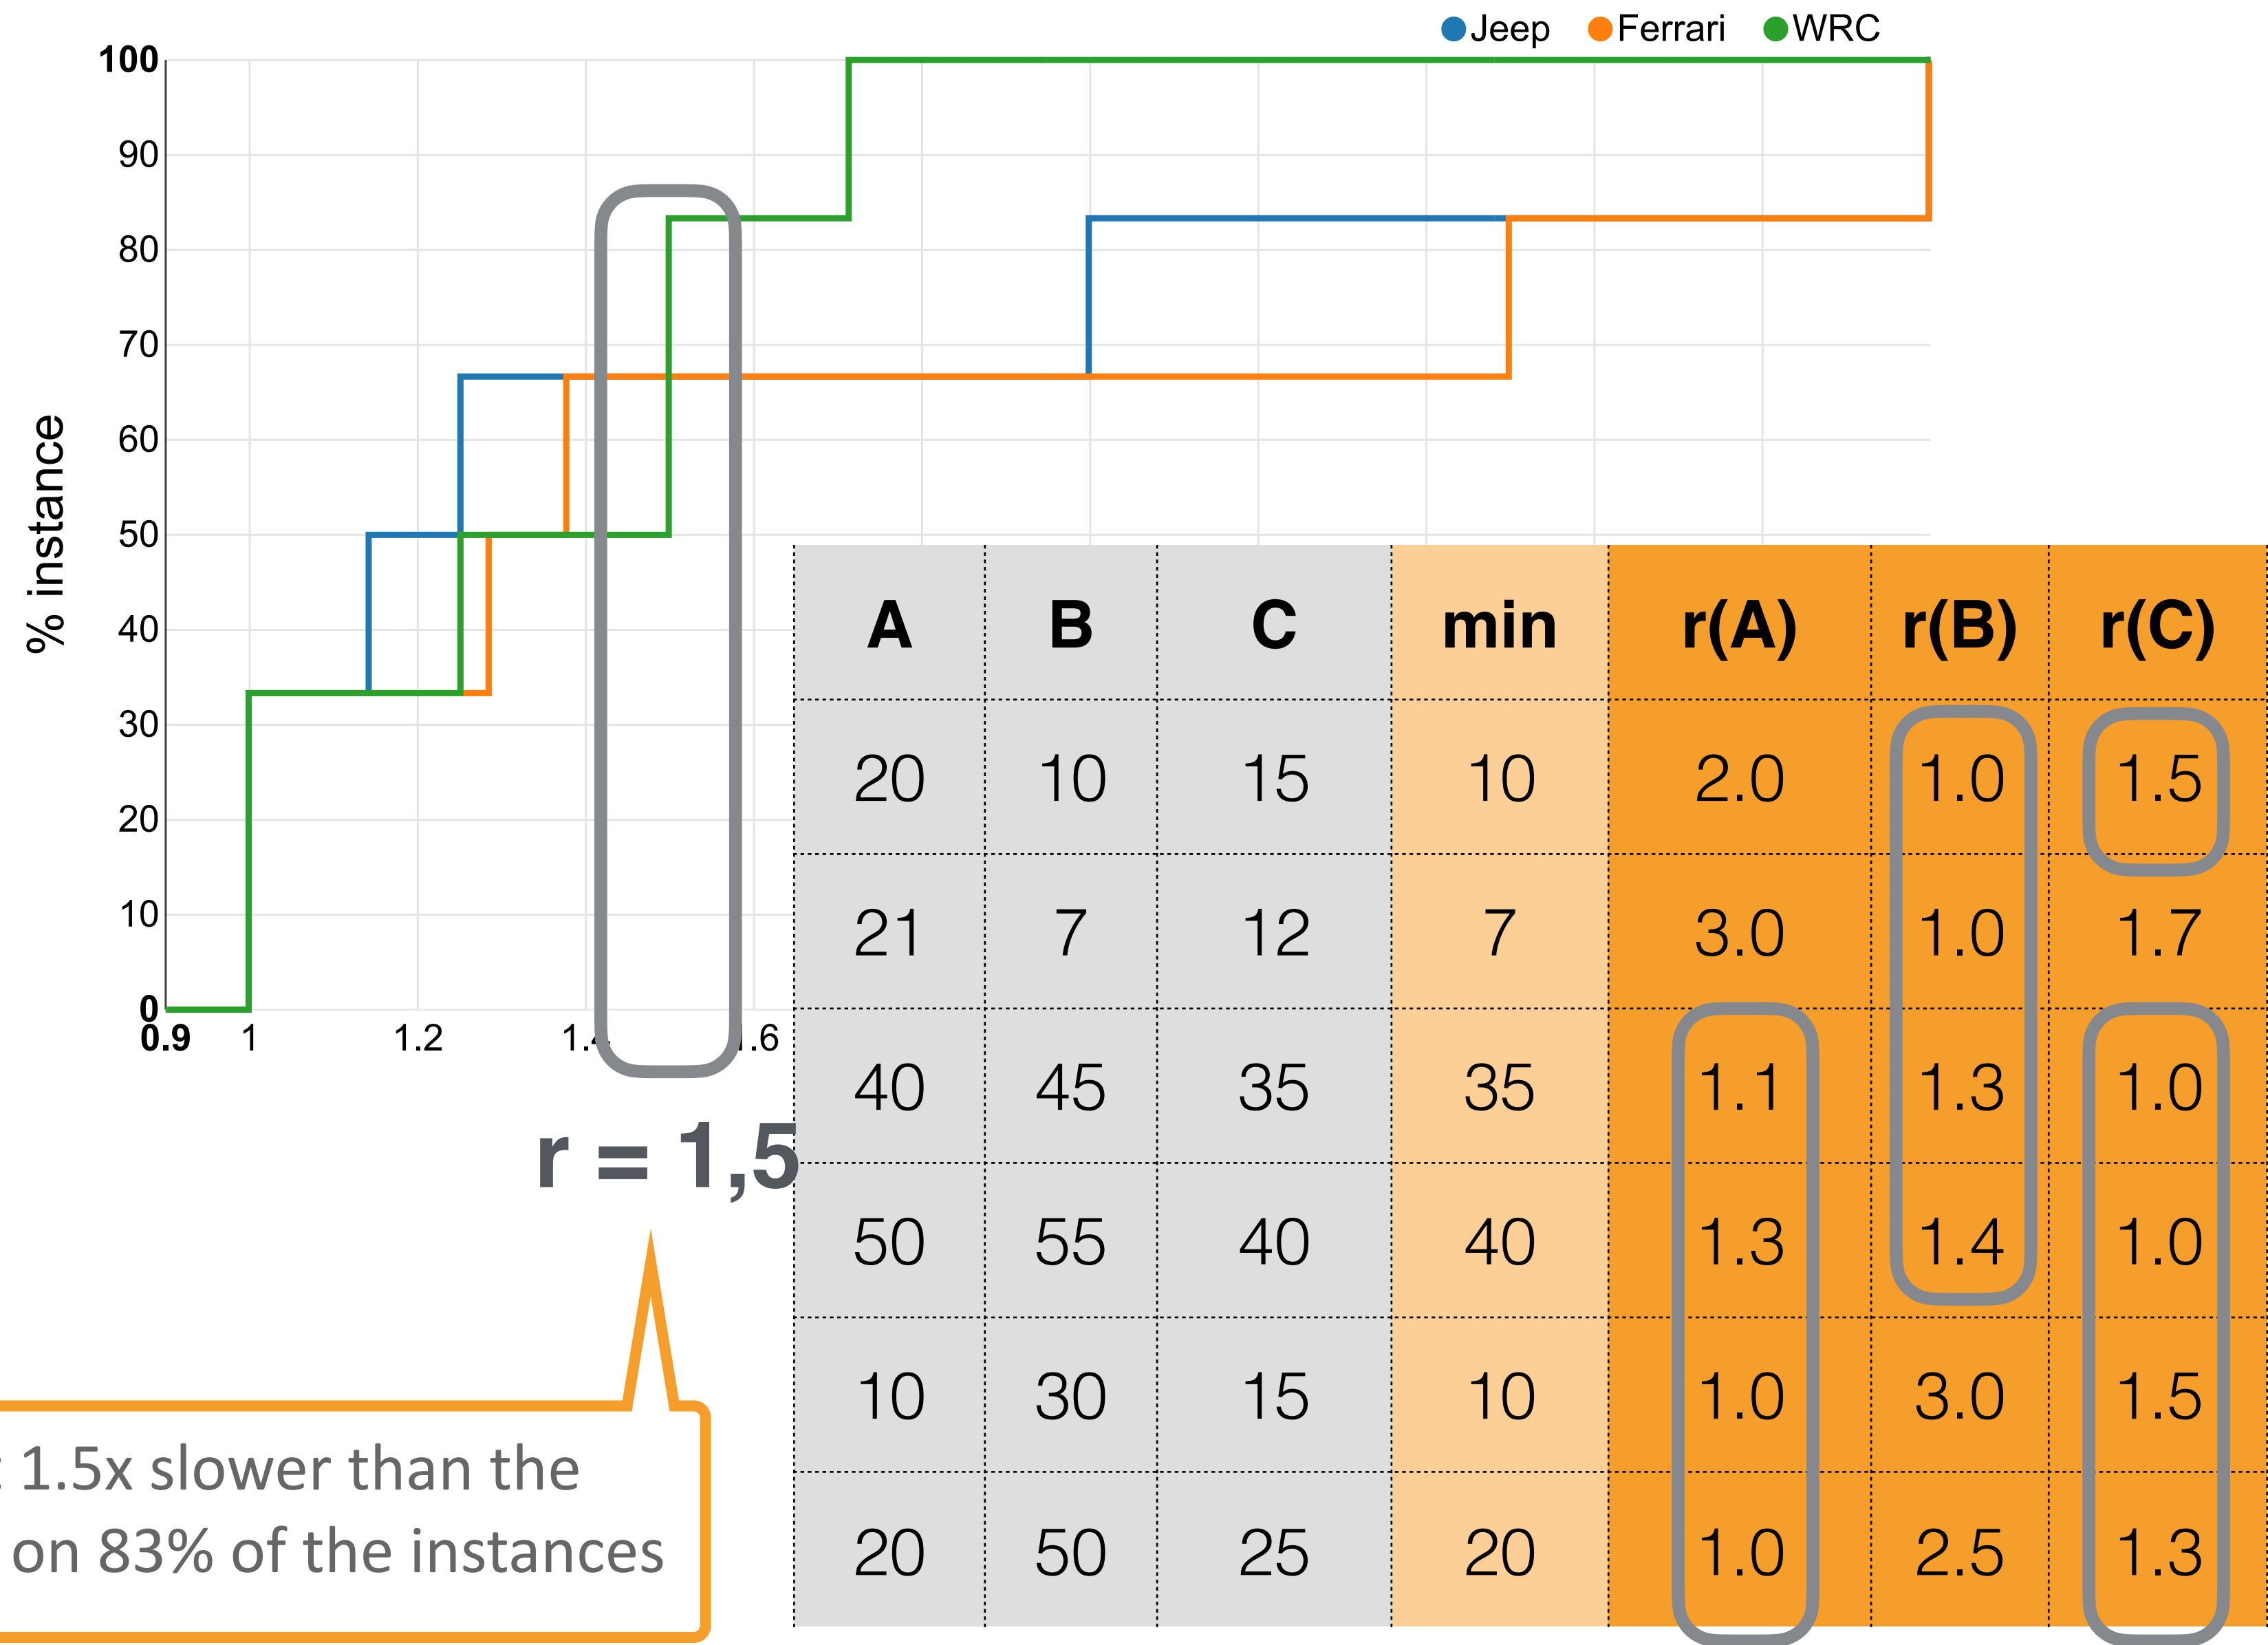
r = 1

A,B,C are the best on 33% of the instances

Performances Profiles Interpretation




Performances Profiles Interpretation



Performance Profile Interpretation: Real Example

LOQO has the highest probability to have the smaller run time




LANCELOT is more robust, it can solve more instances than any other solvers if you allow it to be within a factor ≥ 4 of the best run time.

SNOPT is still raising, would be interesting to see what happens further \Rightarrow log2 scale

Linear Scale

Fig. 4.1. Performance profile on $[0, 10]$

Performance Profile Interpretation: Real Example



now it is clear that if we allow larger computational time, SNOPT is more robust in the sense that it can solve more instances

r_M is our timeout or a measure larger than our timeout

Log Scale

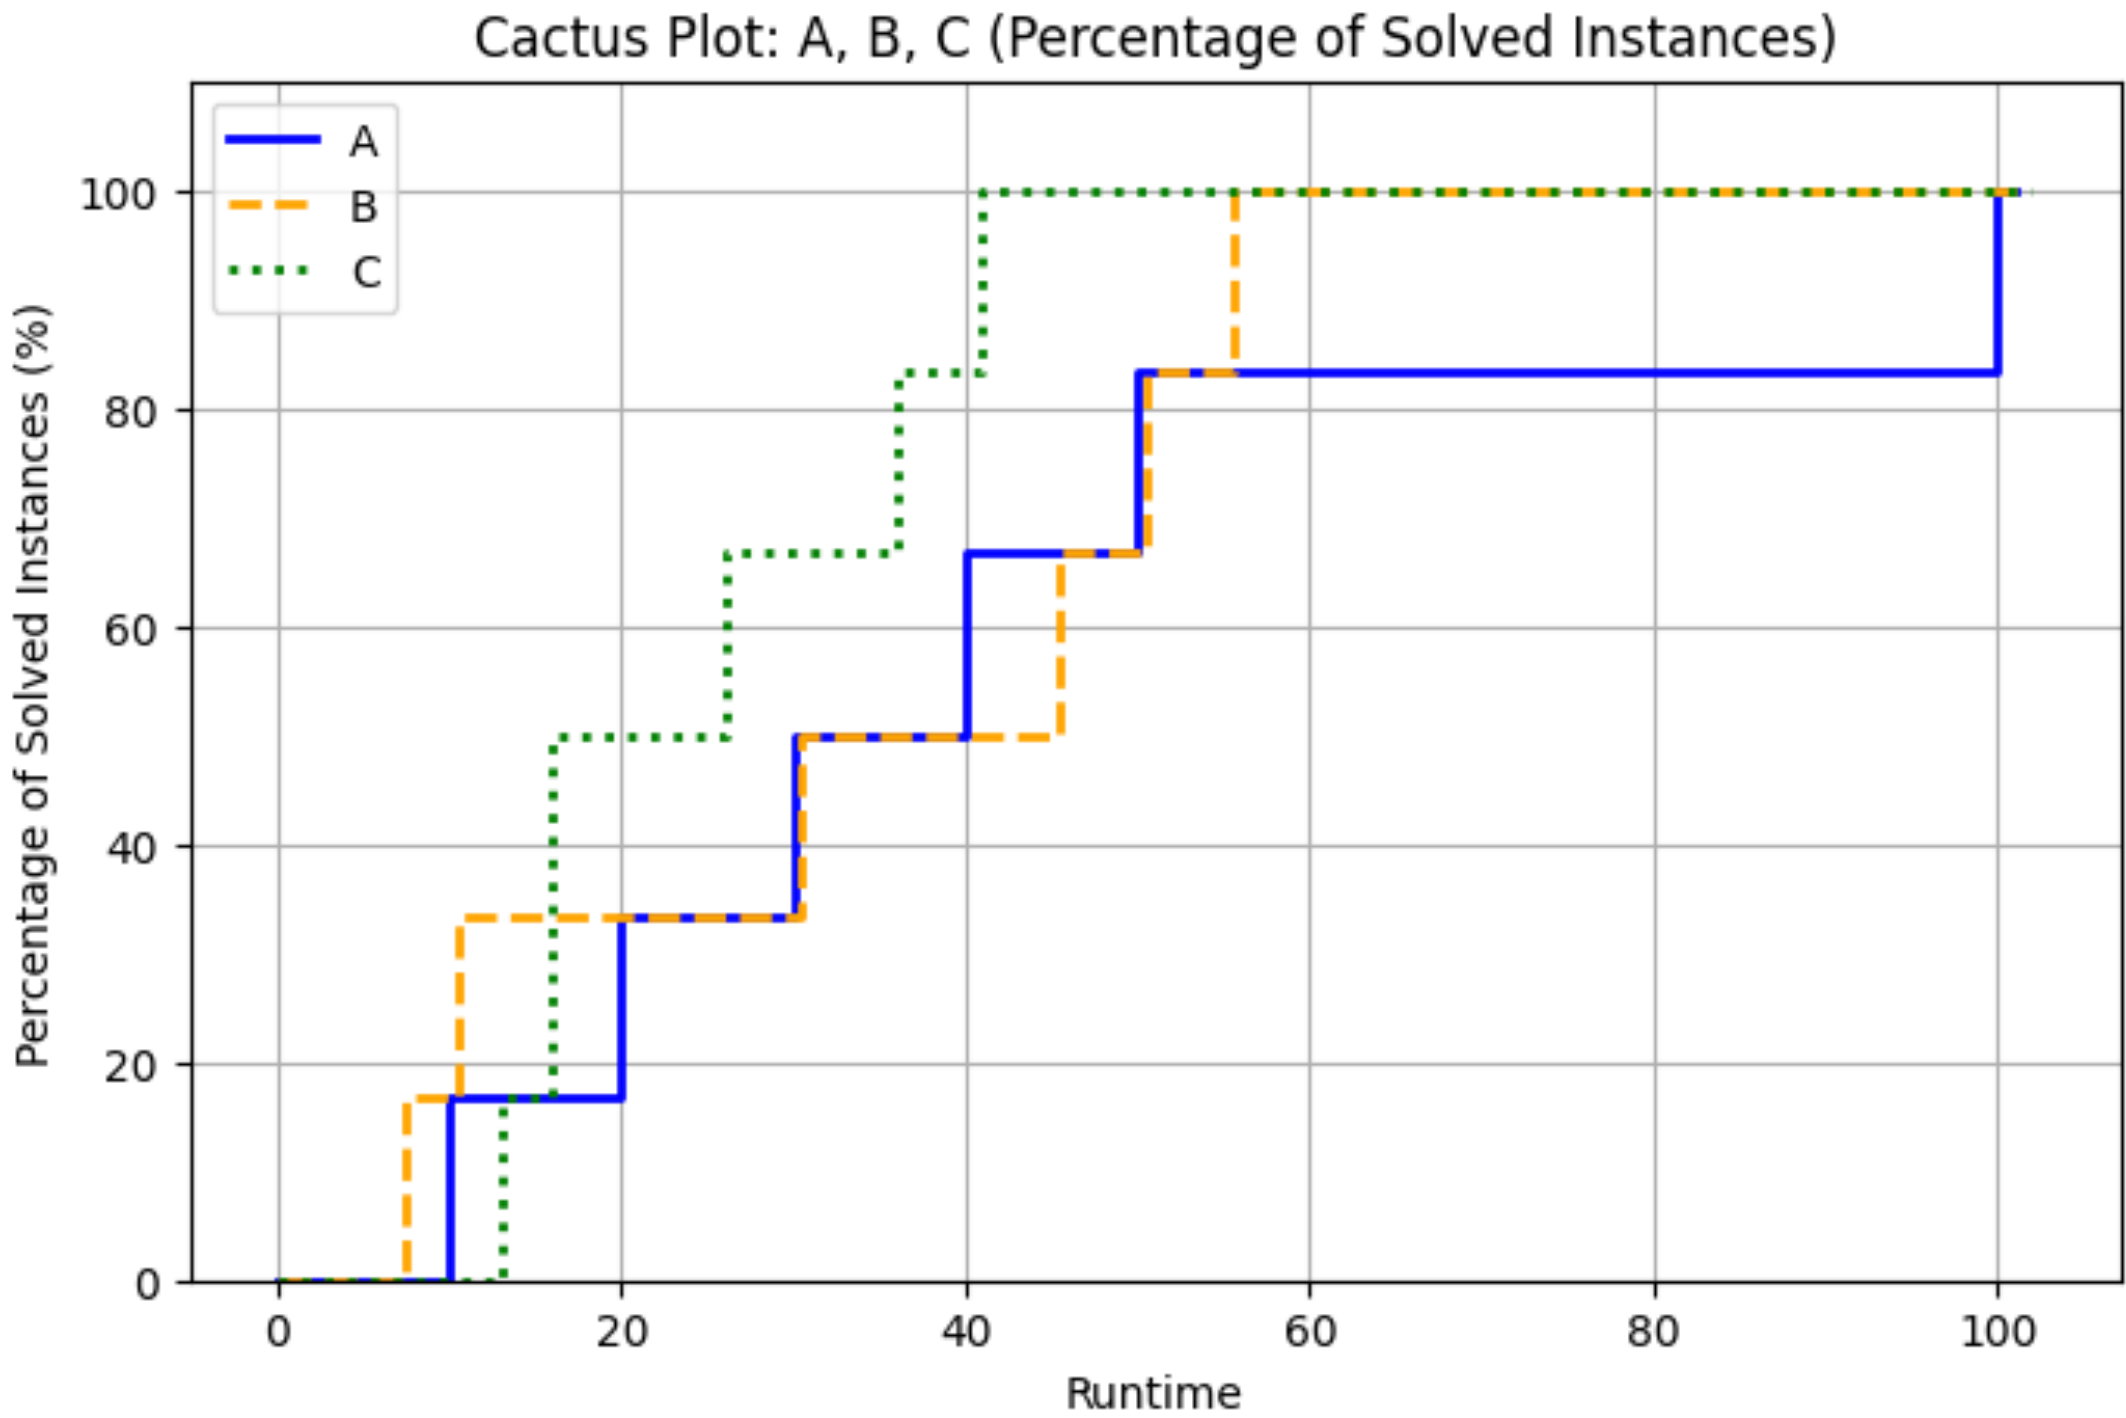
Fig. 4.3. Performance profile in a log₂ scale

Dolan, E. D., & Moré, J. J. (2002). Benchmarking optimization software with performance profiles.

Cactus Plot vs Performance Profiles

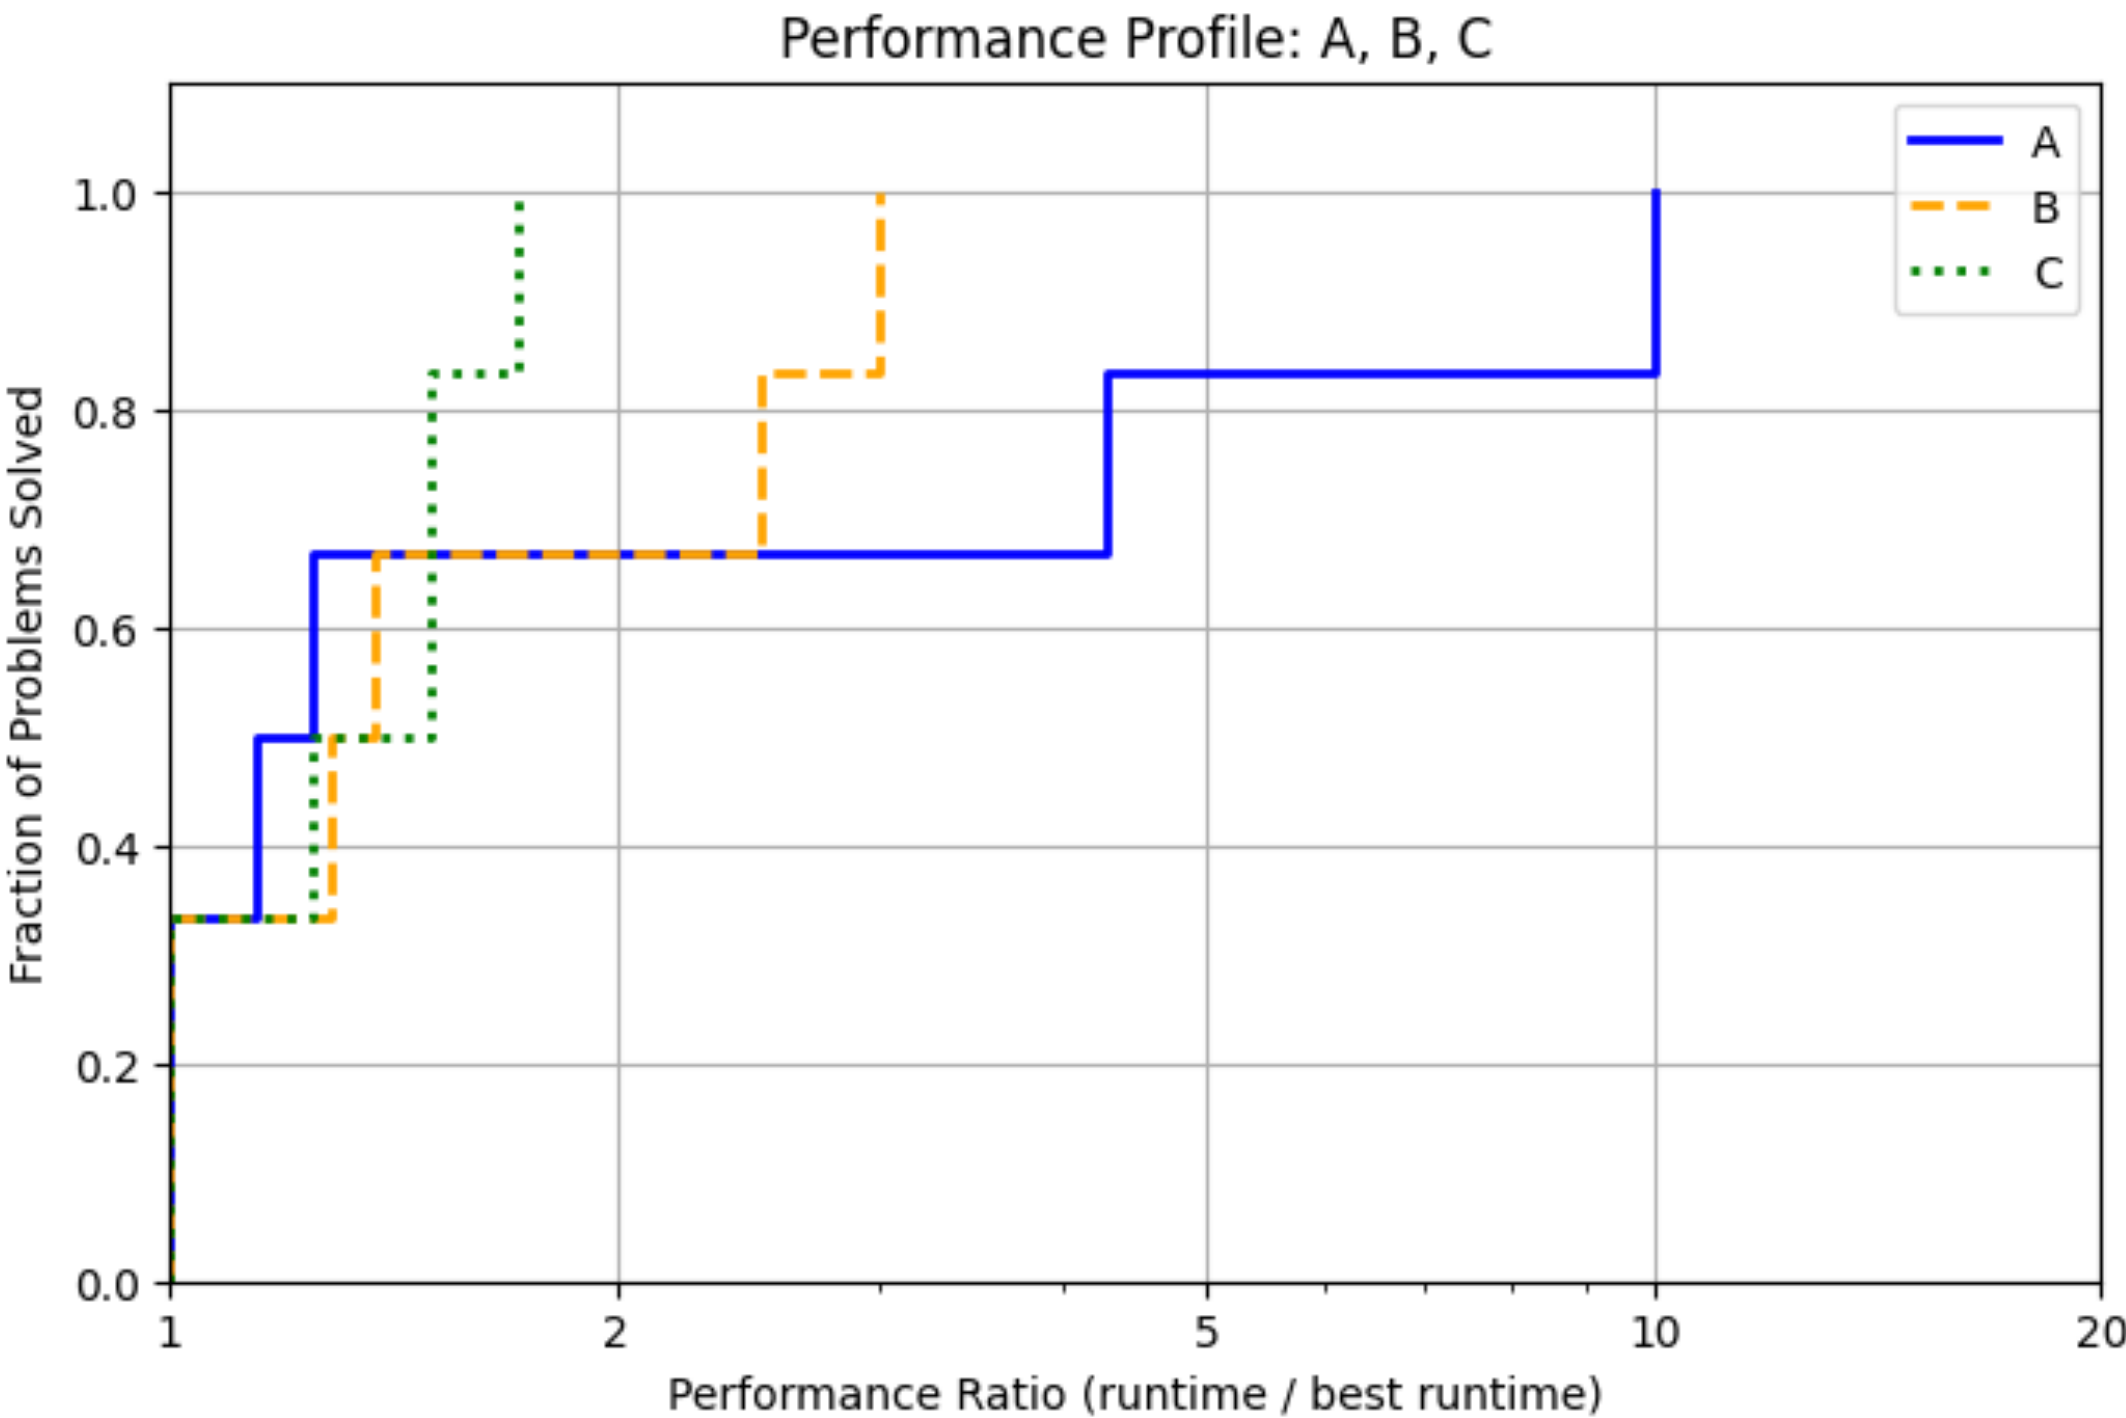
Cactus Plot

$$\kappa_s(\tau) = \frac{1}{n_p} \text{size}\{p \in \mathcal{P} : t_{p,s} \leq \tau\}$$



Performance Profile

$$\rho_s(\tau) = \frac{1}{n_p} \text{size}\{p \in \mathcal{P} : r_{p,s} \leq \tau\}$$




Optimality GAP (Minimization Problem)

- z^* is the optimal objective value of the problem.
- \bar{z} is the best solution found so far

$$\gamma = \frac{|\bar{z} - z^*|}{|\bar{z}|}$$

- Note that if the optimal objective value is not known, a lower bound \underline{z} can be used instead of z^* to compute the optimality gap.
- This lower bound can be provided by the solver itself (best lower bound of open-nodes). MIP solvers report the gap in %

Optimality Gap Decrease: Real Example (Gurobi)

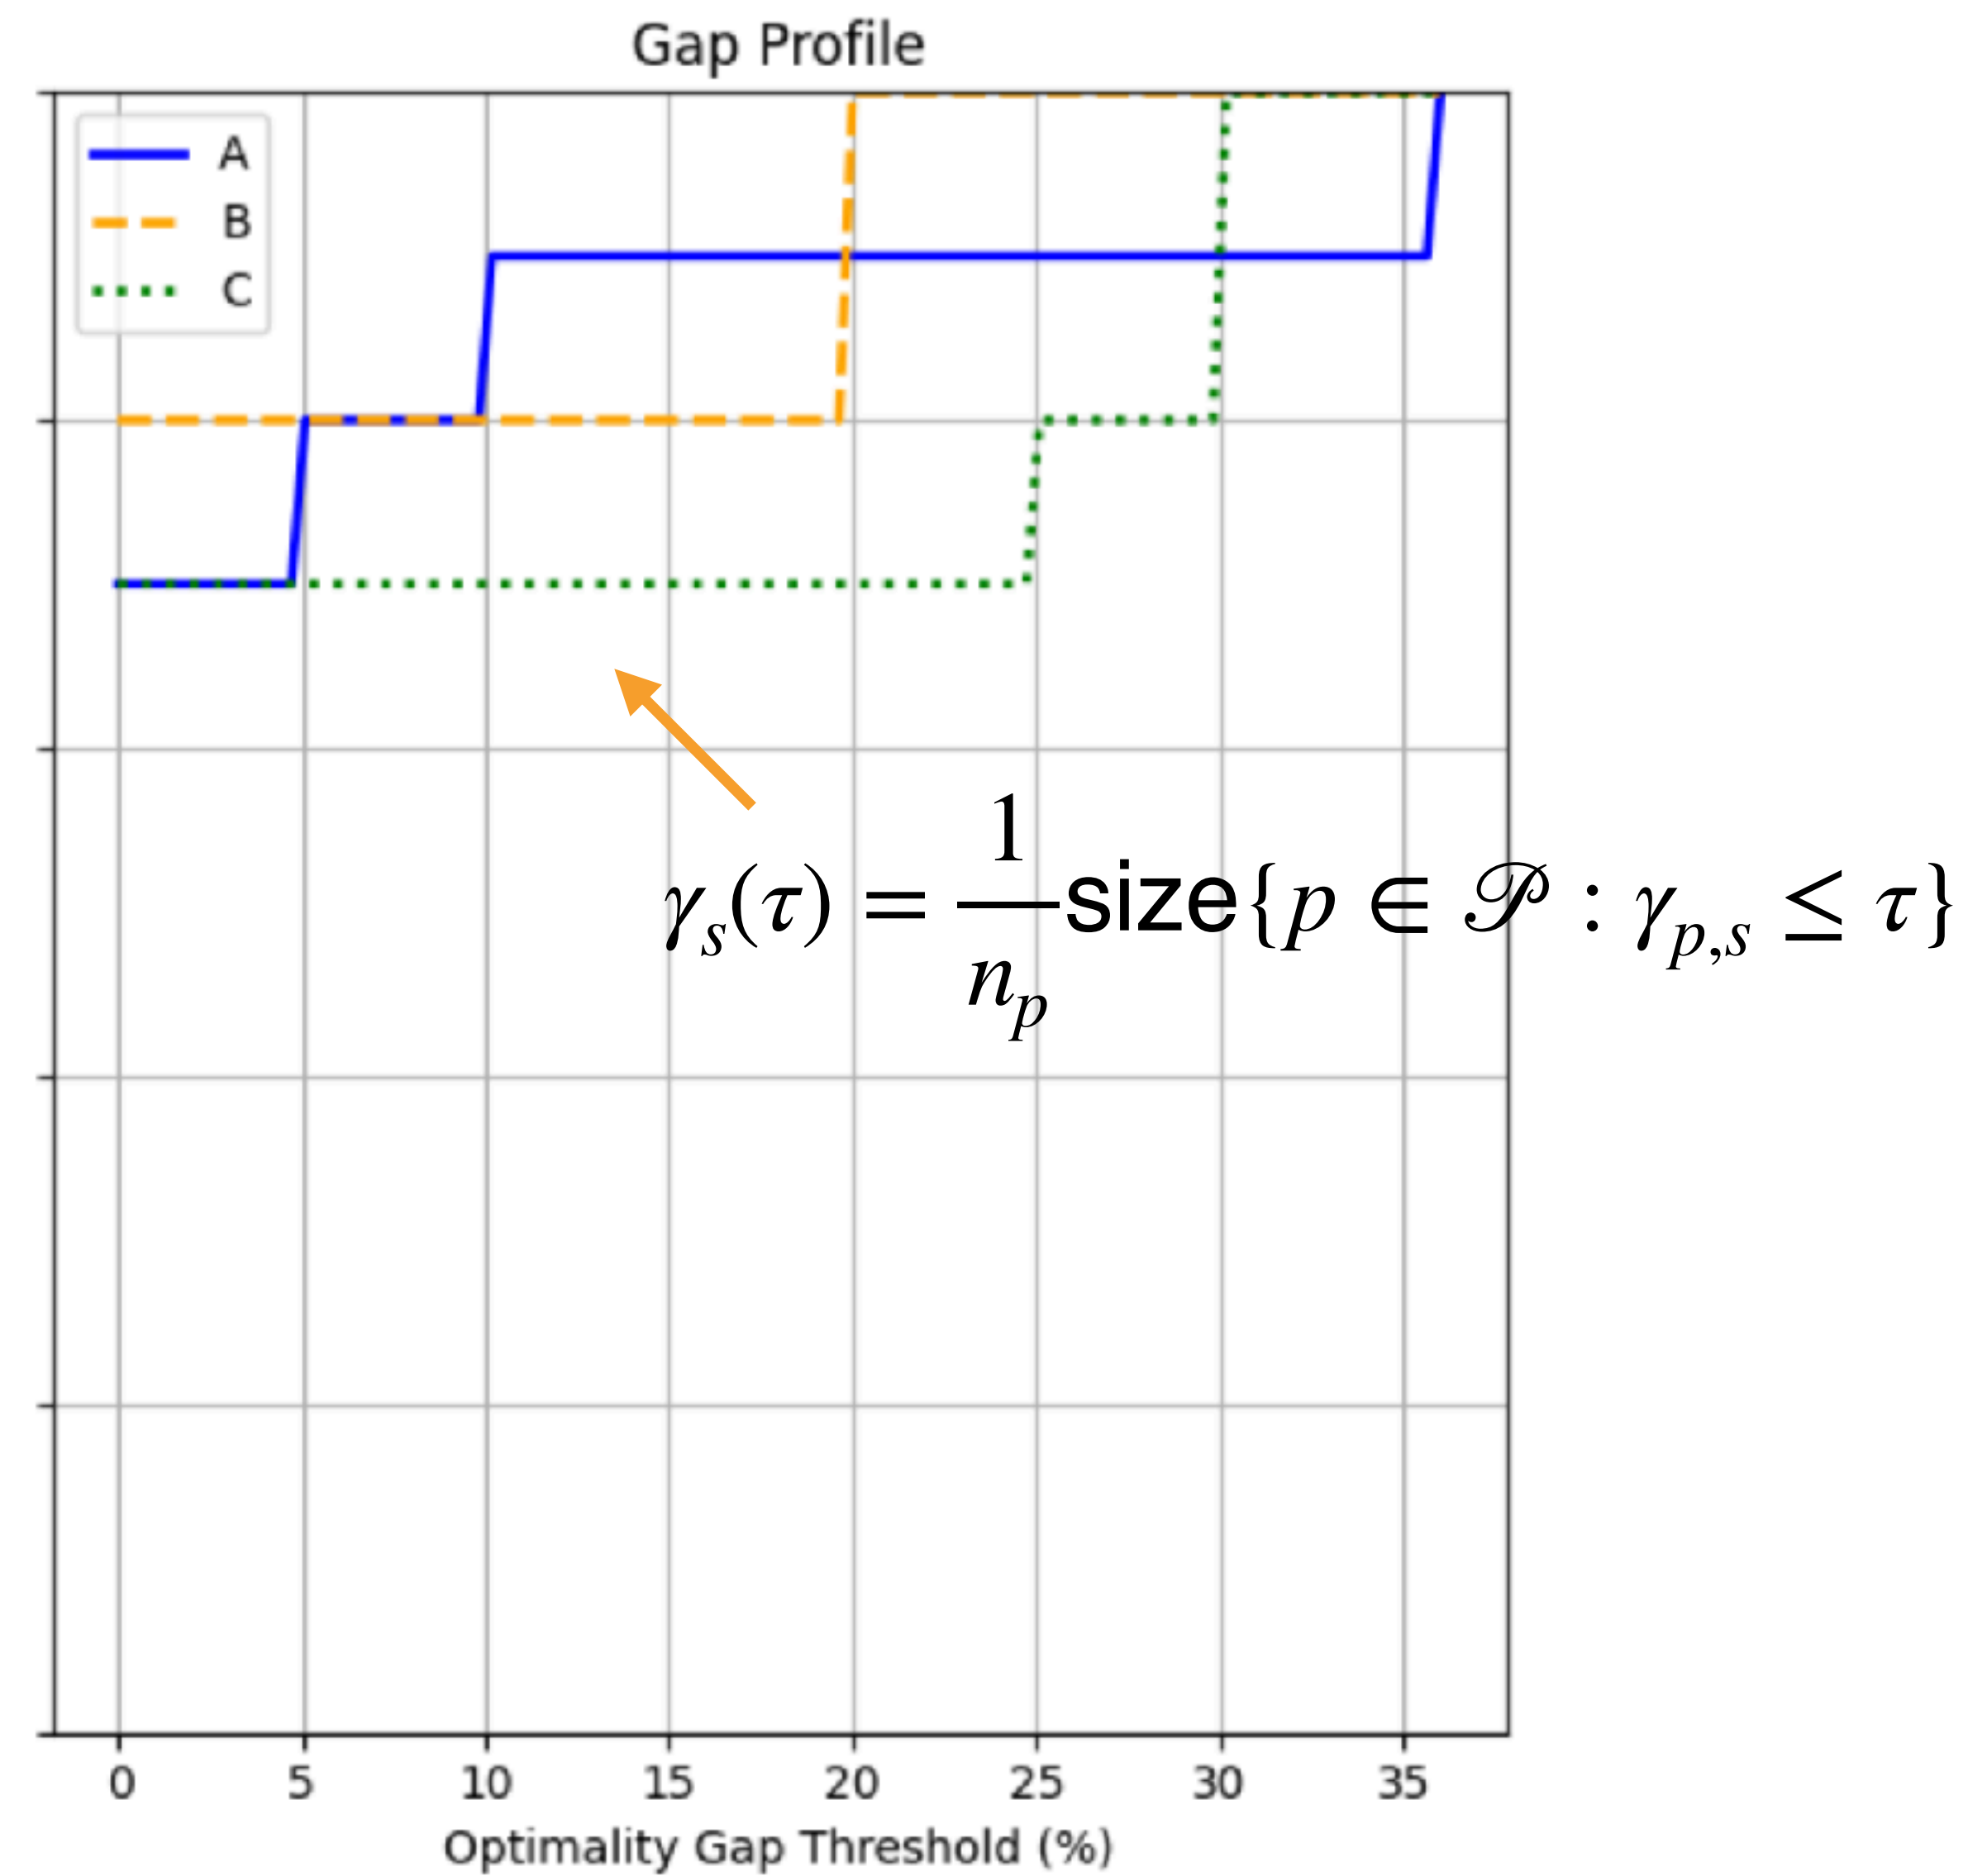



<https://support.gurobi.com/Hc/En-Us/Articles/8265539575953-What-Is-the-MIPGap>

Optimization: Objective Value


- Time to prove optimality + GAP at the time out

Instance	A		B		C	
	Time	γ	Time	γ	Time	γ
1	TO	5	10	0	15	0
2	30	0	7	0	12	0
3	40	0	45	0	35	0
4	TO	36	55	0	TO	30
5	10	0	TO	20	15	0
6	20	0	50	0	25	0
7	TO	10	45	0	TO	25
8	50	0	55	0	TO	30
9	10	0	TO	20	15	0
10	20	0	50	0	25	0



Anytime Algorithm

- An algorithm has good anytime behavior when it is able to find high-quality solutions, even when the search is stopped before completion.



What algorithm do you prefer ?

Primal Gap Revisited

- Let \bar{z} be a solution for minimization problem, and z^* be an optimal (or best known) solution for that problem.
- We define the primal gap $\gamma \in [0,1]$ of \bar{z} as follows:

$$\gamma(\bar{z}) := \begin{cases} 0, & \text{if } |z^*| = |\bar{z}| = 0, \\ 1, & \text{if } z^* \cdot \bar{z} < 0, \\ \frac{|z^* - \bar{z}|}{\max\{|z^*|, |\bar{z}|\}}, & \text{else.} \end{cases}$$

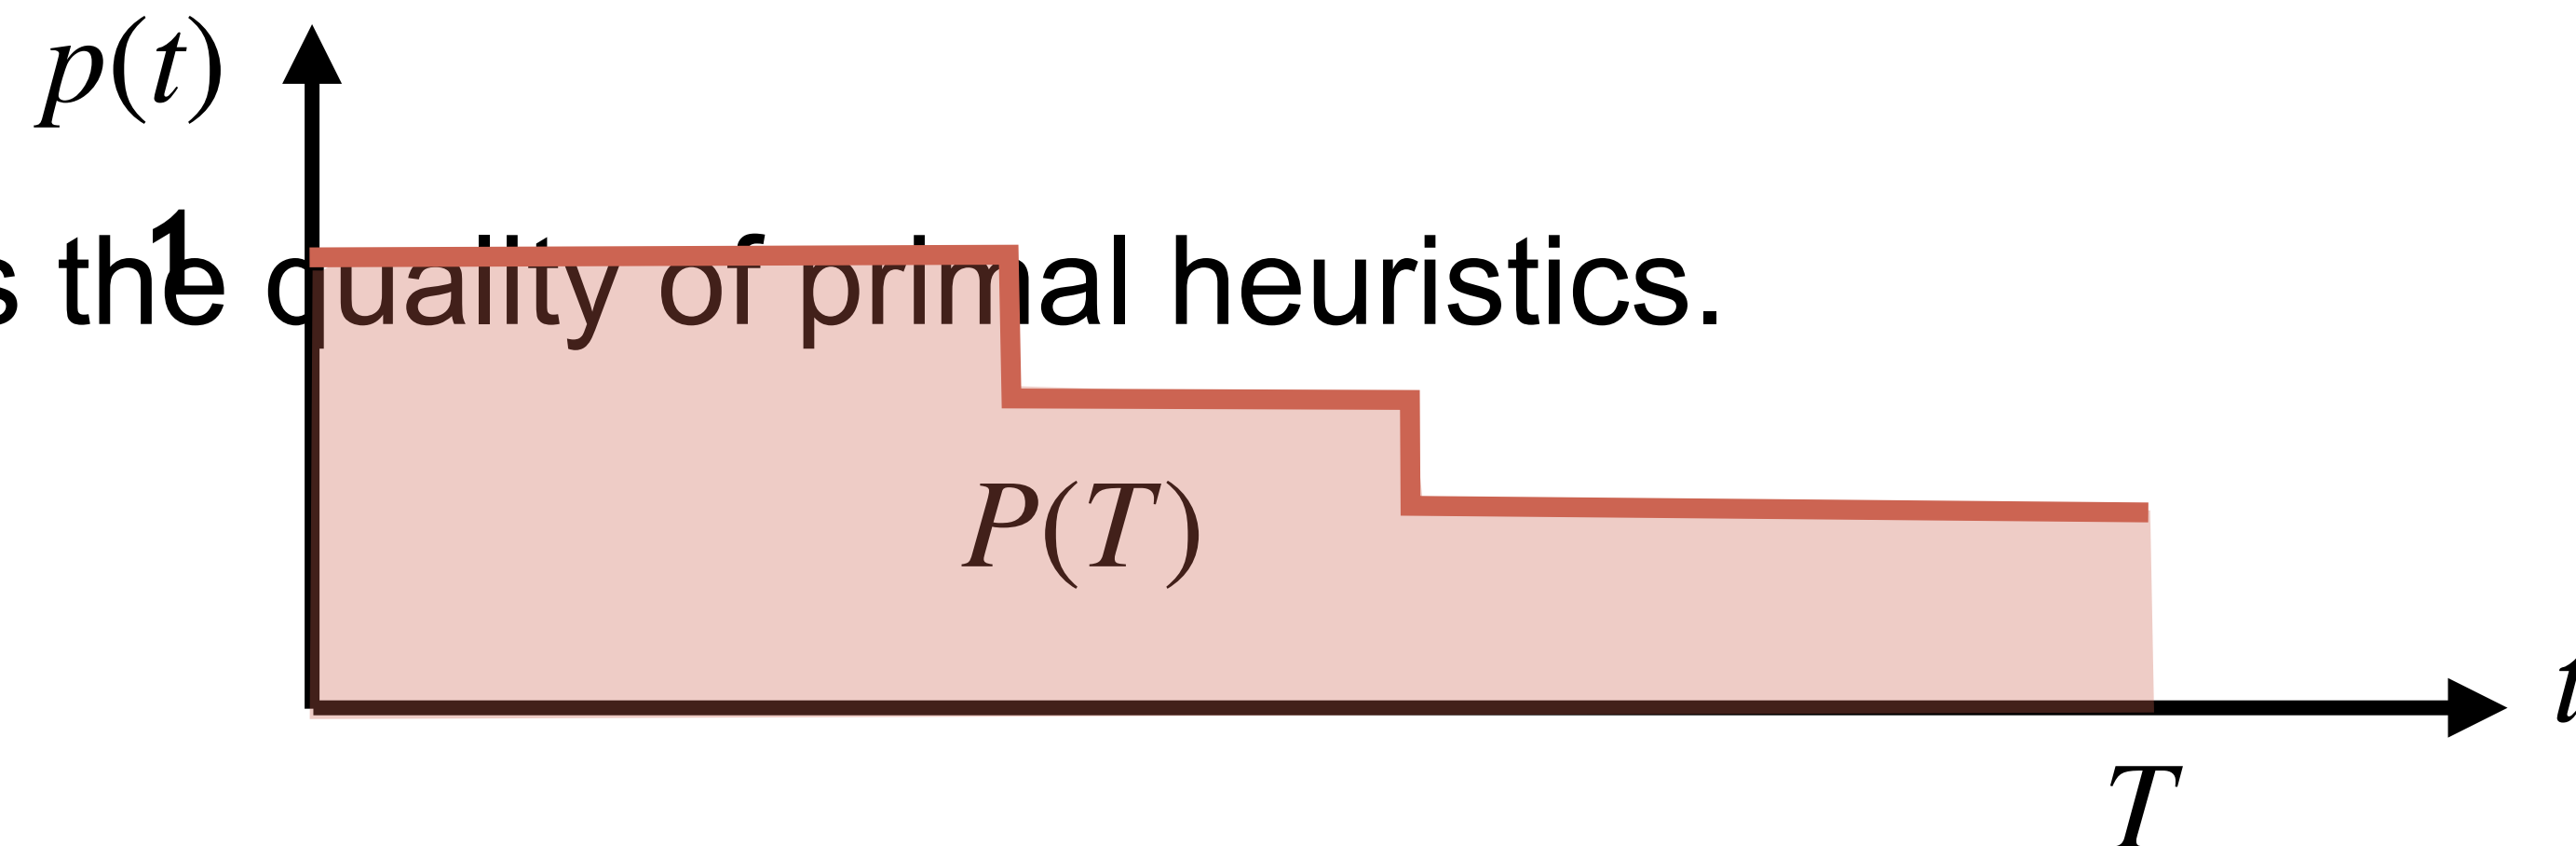
- Now, assume that we have available the objective function values of intermediate incumbent solutions and the points in time when they have been found, for a given a solver, a certain problem instance, and a fixed computational environment.
- Let $t_{\max} \in \mathbb{R}_{\geq 0}$ be a limit on the solution time for the solver.
- Its primal gap function $p : [0, t_{\max}] \mapsto [0, 1]$ is defined as follows:
$$p(t) := \begin{cases} 1, & \text{if no incumbent until point } t, \\ \gamma(\bar{z}(t)), & \text{with } \bar{z}(t) \text{ being the incumbent solution} \\ & \text{at point } t, \text{ else.} \end{cases}$$

Primal Integral

- Let $T \in [0, t_{\max}]$, and let $t_i \in [0, T]$ for $i \in \{1, \dots, I-1\}$ be the points in time when a new incumbent solution is found, $t_0 = 0$, $t_I = T$.
- We define the primal integral $P(T)$ of a run as follows:

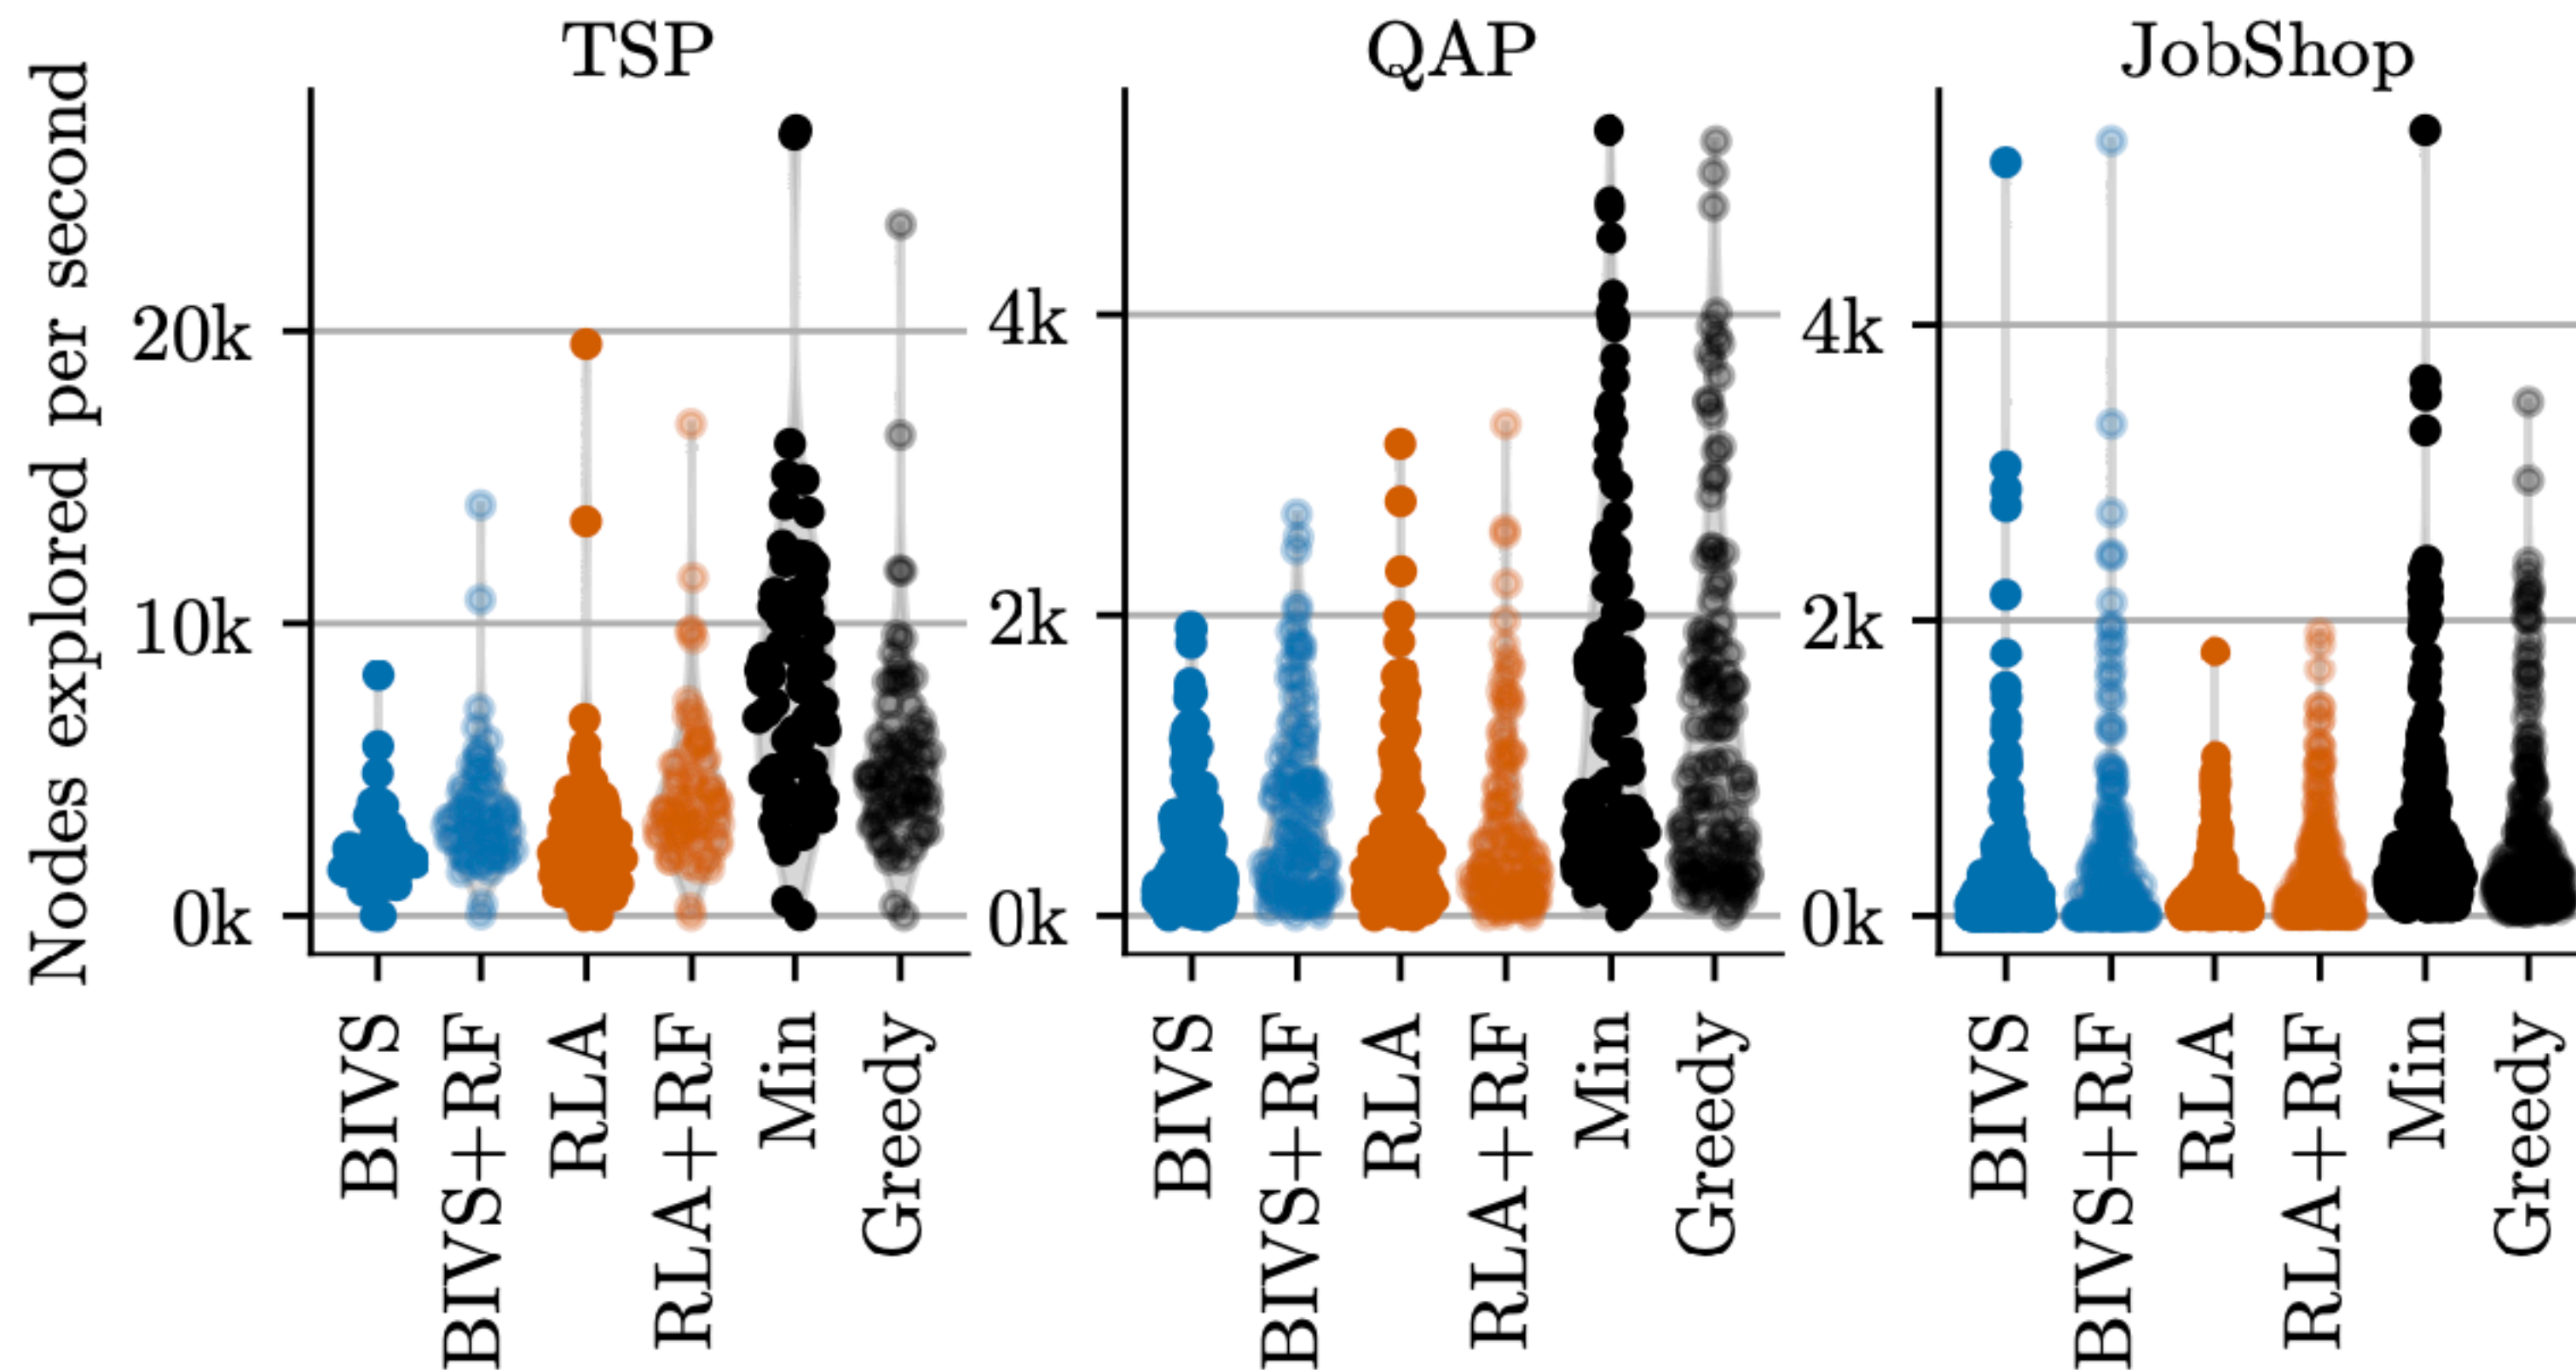
$$P(T) := \sum_{i=1}^I p(t_{i-1}) \cdot (t_i - t_{i-1})$$

- $P(t_{\max})$ measures the quality of primal heuristics.



Other interesting plot: Scatter plot Algo A vs Algo B

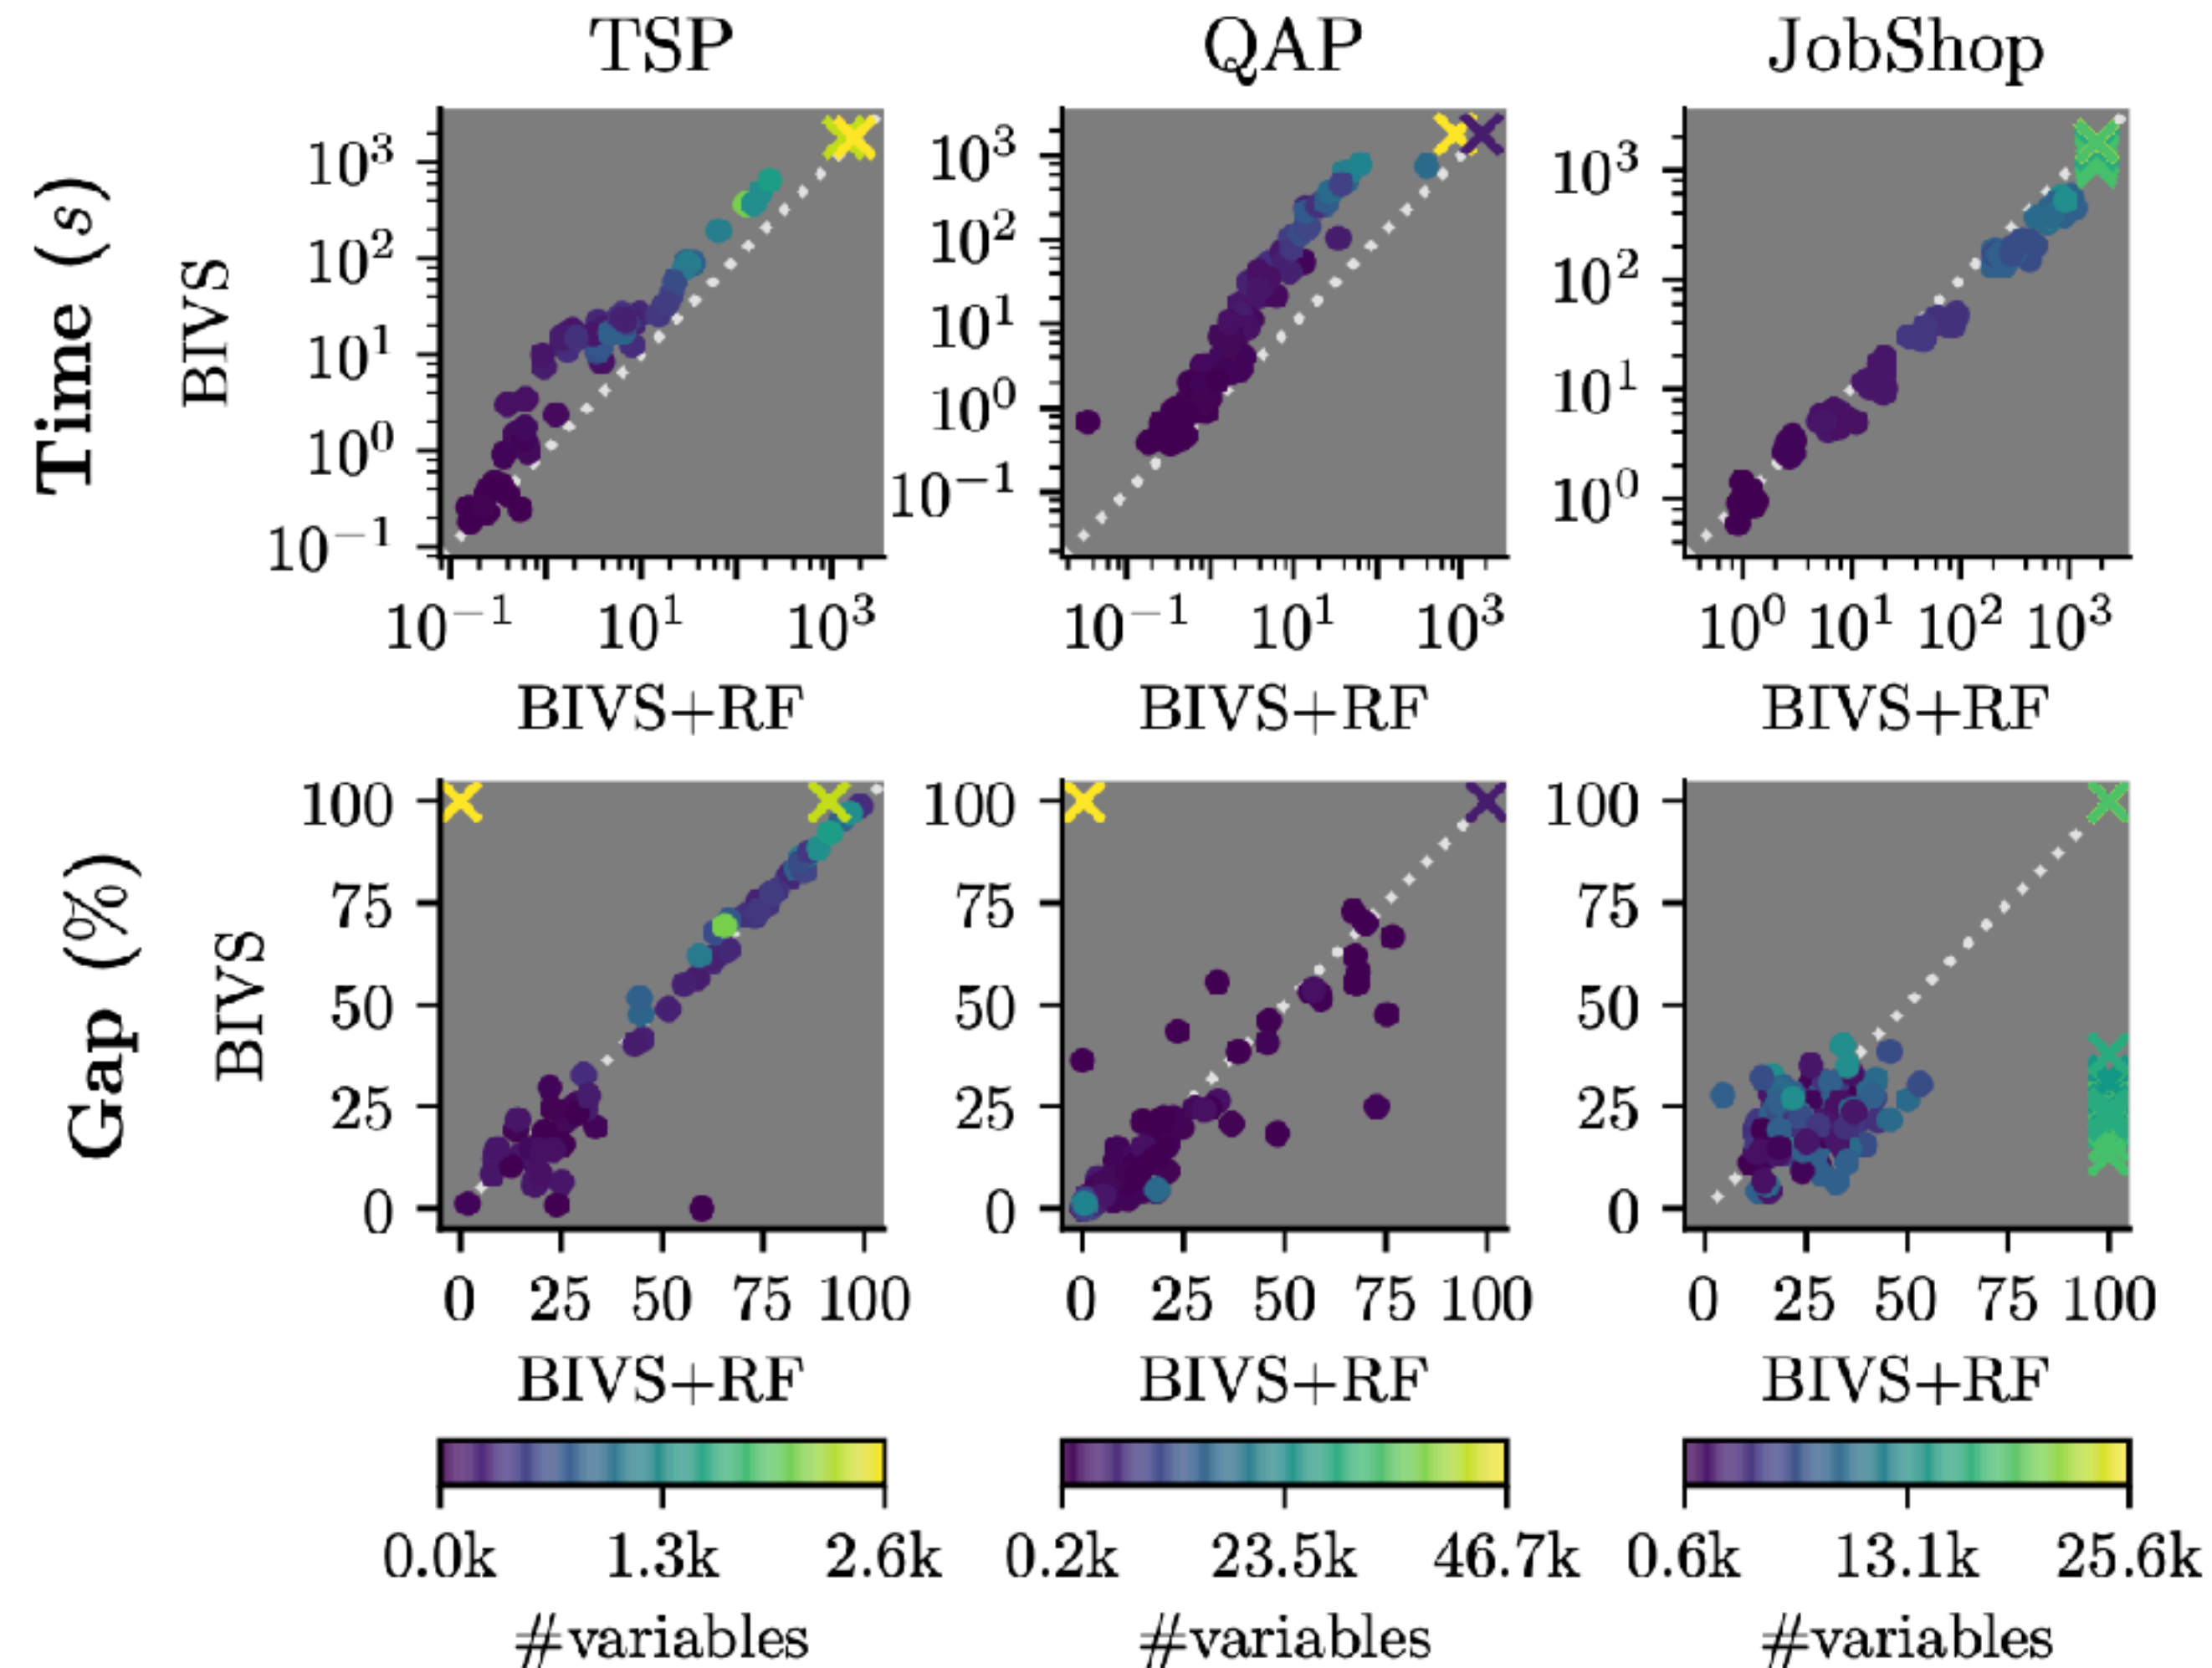
- SinaPlot (more info than box-plot that gives nothing about distribution)



N. Sidiropoulos, S. H. Sohi, T. L. Pedersen, B. T. Porse, O. Winther, N. Rapin, and F. O. Bagger. "SinaPlot: an enhanced chart for simple and truthful representation of single observations over multiple classes". In: *Journal of Computational and Graphical Statistics* 27.3 (2018), pp. 673–676.

Other interesting plot: Scatter plot Algo A vs Algo B

To compare time, backtracks, gaps, etc



⚠ Arithmetic Mean of Normalized Results = 💀

- Assume that for each instance, we have measured the time required to reach optimality for algorithms R,M,Z
- Normalize the results wrt first algorithm (R) and then look at the arithmetic mean of the normalized results to compare.

Benchmark	algorithm		
	R	M	Z
E	417 (1.00)	244 (0.59)	134 (0.32)
F	83 (1.00)	70 (0.84)	70 (0.85)
H	66 (1.00)	153 (2.32)	135 (2.05)
I	39,449 (1.00)	33,527 (0.85)	66,000 (1.67)
K	772 (1.00)	368 (0.48)	369 (0.45)
Arithmetic mean	(1.00)	(1.01)	(1.07)

R is the winner

M is 1% slower than R ?

Arithmetic mean, cont

- Now normalize wrt to algorithm M

Benchmark	algorithm			problem?
	R	M	Z	
E	417 (1.71)	244 (1.00)	134 (0.55)	Its gives
F	83 (1.19)	70 (1.00)	70 (1.00)	
H	66 (0.43)	153 (1.00)	135 (0.88)	
I	39,449 (1.18)	33,527 (1.00)	66,000 (1.97)	
K	772 (2.10)	368 (1.00)	369 (1.00)	
Arithmetic mean	(1.32)	(1.00)	(1.08)	

R is 32% slower than R ?

M is the winner

Solution: Geometric Mean

- With geometric mean, whatever column you choose for normalization, you come to the same winner
- Geometric Mean is the only correct way to aggregate normalized results.
- *HOW NOT TO LIE WITH STATISTICS: THE CORRECT WAY TO SUMMARIZE BENCHMARK RESULTS* (Fleming & Wallace, 1986)

geometric mean = $\left(\prod_{n=1}^k x_n\right)^{\frac{1}{k}}$

algorithm

Benchmark	R	M	Z
E	417 (1.00)	244 (0.59)	134 (0.32)
F	83 (1.00)	70 (0.84)	70 (0.85)
H	66 (1.00)	153 (2.32)	135 (2.05)
I	39,449 (1.00)	33,527 (0.85)	66,000 (1.67)
K	772 (1.00)	368 (0.48)	369 (0.45)
Geometric mean	(1.00)	(0.86)	(0.84)

algorithm

Benchmark	R	M	Z
E	417 (1.71)	244 (1.00)	134 (0.55)
F	83 (1.19)	70 (1.00)	70 (1.00)
H	66 (0.43)	153 (1.00)	135 (0.88)
I	39,449 (1.18)	33,527 (1.00)	66,000 (1.97)
K	772 (2.10)	368 (1.00)	369 (1.00)
Geometric mean	(1.17)	(1.00)	(0.99)